

# Enterprise AI Platform — Architecture Decision Records

June 2026

**Audience:** Architects, Principal Engineers, and anyone who needs to understand why the platform was designed this way. These 20 ADRs form the permanent architectural archive of the project.

## Contents

- [1. Introduction](#)
- [2. Decision Index](#)
- [3. Architecture Timeline](#)
- [4. ADR-001 to ADR-020](#)

## Introduction

Architecture Decision Records capture the context, options, and rationale behind every major engineering decision in this platform. They are written in the present tense ("we decide...") because the decision is always current until a new ADR supersedes it.

Each ADR follows a consistent structure:

- Context** — The situation that motivated the decision
- Decision** — What we decided and why
- Alternatives Considered** — What else we evaluated and why it was rejected
- Consequences** — What became easier or harder as a result
- Trade-offs** — What we accepted and what we sacrificed
- Future Evolution** — How the decision might change

ADRs are never deleted. Superseded decisions are marked "Superseded by ADR-XXX" and kept for historical context. The archive tells the story of the platform's architectural evolution.

## Decision Index

| ADR | Title                     | Status   | Category          | Date    |
|-----|---------------------------|----------|-------------------|---------|
| 001 | Modular Monolith          | Accepted | Architecture      | 2026-06 |
| 002 | Spring Boot 3.3           | Accepted | Technology        | 2026-06 |
| 003 | Provider SPI              | Accepted | AI Infrastructure | 2026-06 |
| 004 | Provider Router           | Accepted | AI Infrastructure | 2026-06 |
| 005 | Prompt Registry           | Accepted | AI Infrastructure | 2026-06 |
| 006 | Model Capability Registry | Accepted | AI Infrastructure | 2026-06 |
| 007 | Semantic Enrichment       | Accepted | Knowledge         | 2026-06 |
| 008 | GraphRAG                  | Accepted | Knowledge         | 2026-06 |
| 009 | Neo4j                     | Accepted | Knowledge         | 2026-06 |
| 010 | Qdrant                    | Accepted | Knowledge         | 2026-06 |

|     |                         |          |               |         |
|-----|-------------------------|----------|---------------|---------|
| 011 | Retrieval Orchestration | Accepted | Knowledge     | 2026-06 |
| 012 | Explainability          | Accepted | Quality       | 2026-06 |
| 013 | Evaluation Engine       | Accepted | Quality       | 2026-06 |
| 014 | Workflow Engine         | Accepted | Extensibility | 2026-06 |
| 015 | AI Observability        | Accepted | Operations    | 2026-06 |
| 016 | Graceful Degradation    | Accepted | Operations    | 2026-06 |
| 017 | DomainConfiguration     | Accepted | Extensibility | 2026-06 |
| 018 | Provenance Graph        | Accepted | Knowledge     | 2026-06 |
| 019 | Testing Strategy        | Accepted | Quality       | 2026-06 |
| 020 | Design Philosophy       | Accepted | Governance    | 2026-06 |

---

## Architecture Timeline

The platform evolved through four major phases, each documented by a set of ADRs:

**Phase 1 — Foundation (ADRs 001, 002, 020)** Module rationalization from 12 to 7 (later 9) modules. Spring Boot 3.3 selected as the application framework. Eight governing principles established.

**Phase 2 — AI Infrastructure (ADRs 003-006)** Provider abstraction via SPI. Provider Router for intelligent model selection. Prompt Registry with versioning. Model Capability Registry for capability-based routing.

**Phase 3 — Knowledge & Retrieval (ADRs 007-011, 018)** Semantic enrichment during ingestion. GraphRAG with Neo4j traversal. Qdrant for vector search. Retrieval orchestration with intent-driven strategy selection. Provenance-aware knowledge graph.

**Phase 4 — Quality & Operations (ADRs 012-017, 019)** Explainability by default. Evaluation engine. Workflow engine. AI observability. Graceful degradation. DomainConfiguration SPI. Multi-layer testing strategy.

---

## Cross-References

ADR-020 (Design Philosophy) is the root — it defines the eight principles that govern all other decisions.

ADR-003 (Provider SPI) is the most referenced — ADRs 004, 005, 006, 011, 012, 015, and 016 all depend on it.

ADR-007 (Semantic Enrichment) is the foundation for ADRs 008 (GraphRAG), 009 (Neo4j), and 018 (Provenance Graph).

ADR-019 (Testing Strategy) validates the entire architecture through 157 tests.

For the full architecture narrative, see **Architecture-Handbook.pdf**. For practical development guidance, see **Developer-Guide.pdf**.

---

## ADRs

### ADR-1

### Status

Accepted. Implemented in the Maven module structure at `pom.xml` .

## Context

The Enterprise AI Platform must support multiple AI-powered applications (document intelligence, contract analysis, financial review, compliance, etc.) while remaining deployable as a single process. The platform must demonstrate clean separation of concerns without the operational complexity of microservices.

## Decision

Adopt a **Modular Monolith** architecture with 9 Maven modules:

|                                     |                                                       |
|-------------------------------------|-------------------------------------------------------|
| <code>platform-audit</code>         | – Immutable audit log, correlation IDs                |
| <code>platform-auth</code>          | – JWT authentication, refresh token rotation          |
| <code>platform-document</code>      | – Document lifecycle, ingestion pipeline              |
| <code>platform-search</code>        | – Hybrid search (keyword + vector + graph)            |
| <code>platform-ai</code>            | – RAG orchestration, enrichment, evaluation, registry |
| <code>platform-neo4j</code>         | – Auto-generated knowledge graph with provenance      |
| <code>platform-workspace</code>     | – Multi-phase workflow wizard                         |
| <code>platform-observability</code> | – Micrometer metrics, health indicators               |
| <code>platform-api</code>           | – REST + Thymeleaf controllers, assembly              |

Module boundaries follow **dependency inversion**: higher-level modules depend on lower-level APIs, never the reverse.

`platform-api` is the assembly module that wires everything together.

## Alternatives Considered

- **Microservices**: Rejected. The platform demonstrates architecture without distributed systems complexity. Microservices add network boundaries, eventual consistency challenges, and deployment overhead that don't benefit a reference implementation.
- **Single JAR without modules**: Rejected. A flat structure would not enforce dependency direction or demonstrate separation of concerns.
- **OSGi modules**: Rejected. Adds runtime modularity complexity. Maven's compile-time enforcement is sufficient.

## Consequences

- **Clear dependency direction**: `api` → `ai` → `search` → `document` → `audit` (audit is a leaf dependency)
- **Compile-time enforcement**: Modules cannot accidentally depend on each other
- **Single deployable**: One `platform-api` Spring Boot application with all modules on the classpath
- **Test isolation**: Each module can be tested independently

## Trade-offs

- Module boundaries are enforced only at compile time, not runtime
- Adding a new module requires POM maintenance
- Some modules (audit) are cross-cutting dependencies of many others

## Future Evolution

- Modules may split if responsibilities grow too large (e.g., `platform-ai` could become `platform-ai-core` + `platform-ai-providers` )
- Spring Modulith could be introduced for runtime module verification
- The architecture deliberately supports future extraction of modules into microservices if needed

See also: [\[\[ADR-002\]\]](#), [\[\[ADR-019\]\]](#), [\[\[ADR-020\]\]](#)

---

## ADR-2

### Status

Accepted. Implemented in `pom.xml` at `spring-boot-starter-parent:3.3.5`.

### Context

The platform must be recognizable to experienced Java engineers, leverage mature ecosystems for security, persistence, and observability, and remain maintainable by teams familiar with enterprise Java.

### Decision

Use **Spring Boot 3.3** with Java 21 as the application framework. All modules depend on `spring-boot-starter-parent` BOM for version management. Key Spring integrations:

- `spring-boot-starter-web` — REST controllers
- `spring-boot-starter-security` — JWT authentication via Nimbus JOSE
- `spring-boot-starter-data-jpa` — PostgreSQL/H2 persistence
- `spring-boot-starter-actuator` — Health endpoints, metrics
- `spring-boot-starter-thymeleaf` — Server-rendered UI pages
- `spring-boot-starter-validation` — Request validation
- `@ConfigurationProperties` — Type-safe configuration binding
- `@ConditionalOnProperty` / `@ConditionalOnBean` — Provider activation

### Alternatives Considered

- **Quarkus**: Rejected. Smaller ecosystem, less familiar to enterprise teams, fewer library integrations.
- **Micronaut**: Rejected. Strong compile-time DI but smaller community and fewer reference architectures.
- **Plain Java with manual DI**: Rejected. Would require reimplementing what Spring provides (security, JPA, scheduling, Actuator, property binding).
- **Spring Boot 2.x**: Rejected. Java 21 virtual threads require Spring Boot 3.x.

### Consequences

- **Rich ecosystem**: Spring Security, Spring Data JPA, Spring Actuator all integrated out of the box
- **Virtual threads**: Java 21 + Spring Boot 3.3 enable `spring.threads.virtual.enabled`
- **Conditional beans**: `@ConditionalOnProperty` enables graceful degradation for optional infrastructure
- **Configuration**: `@ConfigurationProperties` with `platform.*` prefix ensures consistent, type-safe config

### Trade-offs

- Startup time is slower than compile-time DI frameworks
- Spring's "magic" can obscure wiring for newcomers (mitigated by explicit `@Bean` definitions in `SearchInfrastructureConfig`)
- Memory footprint is larger than minimalist frameworks

### Future Evolution

- Spring Modulith could add runtime module verification
- Spring AI could replace custom provider HTTP clients when mature
- GraalVM native image compilation could reduce startup time for serverless deployments

See also: [\[\[ADR-001\]\]](#), [\[\[ADR-003\]\]](#), [\[\[ADR-015\]\]](#)

---

# ADR-3

## Status

Accepted. Implemented in `platform-ai/src/main/java/com/cognitera/platform/ai/api/ChatCompletionProvider.java` and related interfaces.

## Context

AI backends (Ollama, OpenAI, Anthropic, Gemini) expose different APIs, authentication schemes, and request/response formats. Business logic must not depend on any specific provider implementation. Adding a new provider should require zero changes to orchestration code.

## Decision

Define **SPI interfaces** for every AI capability. Implementations are conditionally activated based on configuration:

| Interface              | Implementations                                      | Activation             |
|------------------------|------------------------------------------------------|------------------------|
| ChatCompletionProvider | OllamaChatProvider, OpenAiChatProvider               | @ConditionalOnProperty |
| EmbeddingProvider      | OllamaEmbeddingProvider                              | @ConditionalOnProperty |
| RerankingProvider      | OllamaRerankingProvider                              | @ConditionalOnProperty |
| VectorSearchProvider   | QdrantVectorSearchProvider, NoOpVectorSearchProvider | Conditional on host    |
| GraphSearchProvider    | Neo4jGraphSearchAdapter, NoOpGraphSearchProvider     | Conditional on Neo4j   |
| EvaluationService      | DefaultEvaluationService                             | Always active          |

The `ProviderRouter` ( `DefaultProviderRouter` ) selects a provider using a 4-tier strategy:

1. Model name prefix ( `openai:gpt-4o` → `openai` )
2. Capability match (streaming → capable provider)
3. Preferred provider
4. First available fallback

Business services depend on interfaces, never on concrete implementations.

## Alternatives Considered

- **Hardcoded provider selection in business logic:** Rejected. Would couple orchestration to specific providers and make adding new providers expensive.
- **Spring AI abstraction:** Rejected. At implementation time, Spring AI 1.0.0 provided limited control over provider selection and lacked the capability registry concept. The dependency was removed during build audit.
- **Factory pattern without SPI:** Rejected. An SPI with conditional beans is more idiomatic in Spring and supports auto-discovery.

## Consequences

- **Provider-agnostic business logic:** `AiService` depends on `ChatCompletionProvider` , not Ollama or OpenAI
- **Conditional activation:** Providers activate only when their configuration is present
- **No-code provider addition:** Implement `ChatCompletionProvider` + `@Component` + `@ConditionalOnProperty`

- **Graceful degradation:** When no provider is available, `ProviderRouter` throws a clear exception

## Trade-offs

- Each provider implements its own HTTP client (no shared HTTP infrastructure)
- Provider implementations must handle their own error translation
- Capability discovery is static (seeded in `DefaultModelCapabilityRegistry`) rather than dynamic

## Future Evolution

- Dynamic capability discovery via provider API introspection
- Shared HTTP client infrastructure with configurable retry/timeout
- Provider cost and latency metadata for intelligent routing

See also: [\[\[ADR-004\]\]](#), [\[\[ADR-006\]\]](#), [\[\[ADR-016\]\]](#)

---

## ADR-4

### Status

Accepted. Implemented in `platform-ai/src/main/java/com/cognitera/platform/ai/application/DefaultProviderRouter.java`.

### Context

With multiple AI providers available (Ollama local, OpenAI cloud), the platform must select the appropriate provider for each inference request. Selection criteria include model name, required capabilities, user preference, and provider availability.

### Decision

Implement a **Provider Router** (`ProviderRouter` interface) with a deterministic 4-tier selection strategy:

1. **Model prefix routing:** `openai:gpt-4o` routes to the OpenAI provider, `ollama:qwen2.5` routes to Ollama
2. **Capability-based selection:** If a specific capability is requested (e.g., `STREAMING`), the router selects a provider whose models support that capability, as defined in the `ModelCapabilityRegistry`
3. **Preferred provider:** If the caller specifies a preferred provider, it is selected if available
4. **First available fallback:** The first provider reporting `isAvailable() == true` is selected

If no provider is available, the router throws `IllegalStateException` with a clear message.

Business services call `router.routeChat(InferenceRequest)` — they request capabilities, not specific providers.

### Alternatives Considered

- **Direct injection of a single provider:** Rejected. Would make multi-provider setups impossible and require code changes to switch providers.
- **Random/round-robin selection:** Rejected. Non-deterministic behavior is harder to debug and audit.
- **Cost/latency-aware routing:** Deferred to future evolution. The router architecture supports adding routing dimensions.

### Consequences

- **Deterministic routing:** Same input always produces same routing decision
- **Auditable:** Every routing decision is logged (`log.debug("Routed '{}' to provider '{}'` by model prefix", ...) )

- **Extensible:** Adding a routing dimension means adding a strategy tier, not rewriting the router
- **Clear failure mode:** Unavailable providers are skipped with explicit fallback

### Trade-offs

- Static capability registry requires manual updates when new models are added
- No runtime performance tracking for latency/cost-based routing
- Model prefix convention ( `provider:model` ) is a convention, not enforced by types

### Future Evolution

- Cost-aware routing using `ModelCapability.estimatedCostPer1kTokens()`
- Latency-aware routing using `ModelCapability.estimatedLatencyMs()`
- Region-aware routing for multi-region deployments
- Dynamic capability discovery from provider APIs

See also: [\[\[ADR-003\]\]](#), [\[\[ADR-005\]\]](#), [\[\[ADR-006\]\]](#)

## ADR-5

### Status

Accepted. Implemented in `platform-ai/src/main/java/com/cogniterra/platform/ai/application/DefaultPromptRegistry.java` .

### Context

AI prompts are a critical platform asset. Without versioning, prompt changes cannot be tracked, audited, or rolled back. Without a registry, prompts are scattered across Java string constants, making them impossible to discover or manage.

### Decision

Implement a **Prompt Registry** ( `PromptRegistry` interface) that stores versioned prompt templates with metadata:

| Feature    | Implementation                                                                                                                                      |
|------------|-----------------------------------------------------------------------------------------------------------------------------------------------------|
| Versioning | <code>PromptTemplate.id + PromptTemplate.version</code> → qualified ID "rag-answer/v1"                                                              |
| Categories | <code>PromptTemplate.Category</code> : RETRIEVAL, SUMMARIZATION, EXTRACTION, CLASSIFICATION, EVALUATION, REASONING, WORKFLOW, GRAPH, SEARCH, SYSTEM |
| Variables  | <code>{{variable}}</code> template substitution via <code>render(Map&lt;String, String&gt;)</code>                                                  |
| Metadata   | <code>expectedOutputType</code> , <code>supportedModels</code> , <code>recommendedTemperature</code> , <code>examples</code>                        |
| Discovery  | <code>findByCategory(Category)</code> , <code>listPromptIds()</code> , <code>getLatest(String)</code>                                               |

The default registry ( `DefaultPromptRegistry` ) seeds 8 prompts across 6 categories. Prompts are registered programmatically; in production, they would be loaded from YAML/JSON resources.

Every inference records which prompt was used ( `RetrievalOrchestrationResult.promptTemplateId()` ), enabling full reproducibility and audit trails.

### Alternatives Considered

- **Hardcoded string constants:** Rejected. No versioning, no discovery, no audit trail.

- **Database-backed prompt store:** Deferred. Adds infrastructure dependency. In-memory registry with resource loading is sufficient for a reference implementation.
- **External prompt management service:** Rejected. Over-engineered for a modular monolith.

### Consequences

- **Reproducibility:** Every inference records `promptTemplateId` + `promptTemplateVersion`
- **Discoverability:** `findByCategory()` enables prompt inventory
- **Safe evolution:** New prompt versions can be registered without deleting old ones
- **Regression testing:** Prompts are identifiable assets that can be tested

### Trade-offs

- In-memory storage means prompts reset on restart (acceptable for a reference implementation)
- No prompt validation at registration time (variables are not checked against template)
- Rendering uses simple string substitution, not a template engine (deliberate simplicity)

### Future Evolution

- YAML/JSON resource loading for external prompt definitions
- Prompt validation: verify all declared variables are used and all template variables are declared
- Prompt A/B testing: serve different versions to different users
- Prompt migration tooling for bulk updates

See also: [\[\[ADR-003\]\]](#), [\[\[ADR-004\]\]](#), [\[\[ADR-012\]\]](#)

## ADR-6

### Status

Accepted. Implemented in `platform-ai/src/main/java/com/cogniterra/platform/ai/application/DefaultModelCapabilityRegistry.java` .

### Context

Different AI models have different capabilities (streaming, vision, JSON mode, tool calling, embeddings). Business logic should not hardcode assumptions like "model X supports JSON." Instead, capabilities should be queryable through a central registry, enabling the `ProviderRouter` to make intelligent selection decisions.

### Decision

Implement a **Model Capability Registry** ( `ModelCapabilityRegistry` interface) that describes every available model's capabilities:

| Capability                                                   | Examples                                    |
|--------------------------------------------------------------|---------------------------------------------|
| <code>supportsStreaming</code>                               | qwen2.5:7b (true), nomic-embed-text (false) |
| <code>supportsVision</code>                                  | gpt-4o (true)                               |
| <code>supportsJson</code> / <code>supportsToolCalling</code> | gpt-4o, llama3.2 (true)                     |
| <code>supportsEmbeddings</code>                              | nomic-embed-text (true)                     |
| <code>supportsStructuredOutput</code>                        | gpt-4o (true)                               |

|                          |                                          |
|--------------------------|------------------------------------------|
| maxContextWindow         | 128K (gpt-4o), 32K (qwen2.5)             |
| maxOutputTokens          | 16384 (gpt-4o)                           |
| estimatedLatencyMs       | 500 (qwen2.5 local), 1200 (gpt-4o cloud) |
| estimatedCostPer1kTokens | \$5.00 (gpt-4o), \$0.00 (local Ollama)   |
| preferredUseCases        | "rag", "analysis", "summarization"       |

The registry supports querying by:

- Exact model name: `get("gpt-4o")`
- Provider: `findByProvider("ollama")`
- Capability: `findByCapability(EMBEDDING)`
- Provider + capability: `findByProviderAndCapability("ollama", STREAMING)`

The `DefaultModelCapabilityRegistry` seeds with 6 known models (4 Ollama, 2 OpenAI). The `ProviderRouter` uses the registry for capability-based routing decisions.

### Alternatives Considered

- **Hardcoded if/else chains:** Rejected. Unmaintainable with growing model lists.
- **Runtime model discovery via API:** Deferred. Provider APIs (Ollama `/api/tags`, OpenAI `/models`) could populate the registry dynamically, but add latency and failure modes.
- **No registry — trust the provider:** Rejected. The platform needs to know capabilities before calling a model (e.g., "does this model support JSON mode?").

### Consequences

- **Capability-aware routing:** `ProviderRouter` selects models based on required capabilities
- **Static seed data:** Registry is populated at startup with known models; new models require code changes
- **Query interface:** Rich query API enables future use cases (e.g., "find the cheapest model that supports vision")

### Trade-offs

- Static data can become stale as providers add new models
- Estimated costs and latencies are approximations, not real-time measurements
- No runtime validation that the model actually supports claimed capabilities

### Future Evolution

- Dynamic discovery from provider `/models` endpoints
- Runtime capability verification (send test prompts to verify claims)
- User-provided model registrations via configuration
- Integration with provider cost APIs for real-time pricing

See also: [\[\[ADR-003\]\]](#), [\[\[ADR-004\]\]](#), [\[\[ADR-005\]\]](#)

---

## ADR-7

### Status


Accepted. Implemented in `platform-ai/src/main/java/com/cogniterra/platform/ai/application/DefaultEnrichmentService.java` and `platform-api/src/main/java/com/cogniterra/platform/api/ingestion/EnrichmentHook.java`.

## Context

Documents contain unstructured information. To enable intelligent retrieval, the platform must extract structured knowledge — entities, concepts, and relationships — automatically during ingestion. Users should never manually populate knowledge bases.

## Decision

Implement an automatic **Semantic Enrichment Engine** that runs as part of the document ingestion pipeline:



The enrichment pipeline uses a dual-mode extraction strategy:

1. **LLM-based:** When a `ChatCompletionProvider` is available, prompts the LLM with `entity-extraction/v1` template for high-quality structured extraction
2. **Regex fallback:** When no LLM is available, uses regex patterns for known entity types (ORGANIZATION, PERSON, DATE, MONEY)

Results are bridged to Neo4j via `EnrichmentHook`, which converts `EnrichmentContext` to `GraphNode` / `GraphRelationship` objects and persists them through `GraphEnrichmentService`.

Entities carry **provenance**: `sourceDocumentId`, `chunkId`, `extractionConfidence`, `extractionTimestamp`, `extractionModel`, `promptVersion`, `provider`.

## Alternatives Considered

- **Manual knowledge entry (old platform-knowledge module):** Rejected. Users should not manually curate knowledge. The document corpus is the knowledge source. The old `platform-knowledge` module was removed in Phase 1.
- **LLM-only enrichment:** Rejected. Would fail when no LLM is available. Regex fallback ensures basic enrichment always works.
- **No enrichment:** Rejected. Without enrichment, retrieval is purely lexical/semantic with no structured understanding.

## Consequences

- **Automatic knowledge extraction:** Entities, concepts, and relationships are extracted during ingestion without user intervention
- **Graceful degradation:** Regex fallback when LLM is unavailable
- **Provenance:** Every graph node links back to its source document and extraction method
- **Neo4j population:** The knowledge graph is auto-generated, never manually curated

## Trade-offs

- Regex patterns are less accurate than LLM extraction
- The regex patterns are English-centric and need extension for multilingual documents
- Enrichment adds latency to the ingestion pipeline

## Future Evolution

- Multilingual entity extraction patterns
- Domain-specific enrichment via `DomainConfiguration`
- Confidence threshold configuration for entity filtering

- Parallel enrichment for large documents
- Integration with external NER services

See also: [[ADR-008]], [[ADR-009]], [[ADR-017]], [[ADR-018]]

---

## ADR-8

### Status

Accepted. Implemented in `Neo4jGraphSearchAdapter` bridging `GraphEnrichmentService` to `GraphSearchProvider`, with integration in `DefaultHybridRetrievalService`.

### Context

Traditional RAG retrieves chunks via keyword and vector similarity. However, semantically related documents may not share keywords or embedding proximity. A knowledge graph connecting entities, concepts, and documents enables traversal-based discovery that complements keyword and vector retrieval.

### Decision

Implement **GraphRAG** — graph-enhanced retrieval — as an optional third retrieval source alongside keyword and vector:

```
SearchQuery
├─ KeywordSearchProvider.search() → keywordResults
├─ VectorSearchProvider.search() → vectorResults
├─ GraphSearchProvider.search() → graphResults (optional)
└─ merge(keyword, vector, graph) → fused candidates
    ├─ combine(): weighted linear fusion ( $k \times 0.40 + v \times 0.40 + c \times 0.20$ )
    └─ combineWithGraph(): graph boost (existing +  $\text{graph} \times 0.15$ )
→ Reranking → LLM
```

Graph results **boost** existing candidates rather than replacing them. The `combineWithGraph()` method adds up to 15% score increase when graph traversal finds related nodes.

Graph retrieval is **optional**: `SearchMode.GRAPH` and `SearchMode.HYBRID_GRAPH` activate it. When Neo4j is unavailable, `NoOpGraphSearchProvider` returns empty results and retrieval continues with keyword + vector.

### Alternatives Considered

- **Graph-only retrieval**: Rejected. Graph traversal is useful for discovery but less precise for exact match queries.
- **Graph as primary source with keyword/vector as secondary**: Rejected. Documents without graph entities would be invisible.
- **No graph retrieval**: Rejected. The enrichment engine already populates Neo4j; not using it for retrieval wastes the enrichment investment.

### Consequences

- **Three-source fusion**: Keyword, vector, and graph results are merged in a single pipeline
- **Graceful degradation**: Graph retrieval is optional; platform works without Neo4j
- **Graph boost**: Related entities and concepts boost document scores by up to 15%
- **Explainability**: Graph participation is tracked in retrieval metadata

### Trade-offs

- Graph search uses simple entity name matching from the query (not embedded query → nearest neighbor in graph embedding space)
- The `Neo4jGraphSearchAdapter` creates synthetic `ChunkReference` objects for graph results, which have lower fidelity than real chunk references
- Graph boost of 15% is fixed, not calibrated per domain

### Future Evolution

- Graph embedding models for semantic graph traversal
- Configurable graph boost factors per domain
- Multi-hop reasoning via graph traversal patterns
- Graph-native reranking using centrality metrics

See also: [\[\[ADR-007\]\]](#), [\[\[ADR-009\]\]](#), [\[\[ADR-011\]\]](#)

---

## ADR-9

### Status

Accepted. Implemented in `platform-neo4j/src/main/java/com/cognitera/platform/neo4j/service/GraphEnrichmentService.java` .

### Context

The platform needs to store structured knowledge extracted from documents — entities, concepts, and their relationships. This data is inherently graph-shaped: entities relate to documents, concepts relate to entities, documents cite other documents. A relational database would require complex recursive CTEs for graph traversal.

### Decision

Use **Neo4j** as the persistence layer for the automatically-generated knowledge graph. Neo4j is **not** a general-purpose database in this architecture — it stores only semantic enrichment output:

| Node Type | Example                                      |
|-----------|----------------------------------------------|
| DOCUMENT  | Uploaded PDF, DOCX, TXT                      |
| ENTITY    | Organization, Person, Technology             |
| CONCEPT   | Temporal concepts, financial amounts, topics |

| Relationship | Example             |
|--------------|---------------------|
| MENTIONS     | Document → Entity   |
| RELATED_TO   | Document → Concept  |
| BELONGS_TO   | Entity → Concept    |
| REFERENCES   | Document → Document |

The graph is **auto-generated during ingestion** — never manually curated. The old `platform-knowledge` module (manual CRUD knowledge base) was removed in Phase 1.

Neo4j is **optional**: the `graph` profile in `docker-compose.yml` and `@ConditionalOnProperty(name = "platform.neo4j.uri")` ensure the platform starts without it.

## Alternatives Considered

- **PostgreSQL with recursive CTEs**: Rejected. Graph traversal queries become complex and slow beyond 2-3 hops. Cypher is purpose-built for graph traversal.
- **Property graph in application memory**: Rejected. Does not persist across restarts; cannot scale beyond small corpora.
- **RDF triple store**: Rejected. Adds complexity (SPARQL, ontology management) without clear benefit over labeled property graphs for this use case.
- **No graph database**: Rejected. The enrichment engine produces graph-shaped data; storing it relationally would violate the "right tool for the data shape" principle.

## Consequences

- **Native graph traversal**: `GraphEnrichmentService.traverse(seedIds, maxDepth)` uses Cypher for multi-hop queries
- **Optional infrastructure**: Platform works without Neo4j
- **Provenance**: Every node carries `NodeProvenance` linking back to source document and extraction method
- **Auto-generated**: Graph population happens during ingestion, not through user interaction

## Trade-offs

- Adds infrastructure dependency when GraphRAG is desired
- Neo4j Community Edition has single-database limitation
- Graph schema is implicit (defined by code) rather than explicit (defined by constraints)

## Future Evolution

- Graph embedding generation in Neo4j (GDS library)
- Cypher query templates for common traversal patterns
- Graph-native reranking using PageRank or centrality algorithms
- Multi-tenancy via Neo4j database-per-tenant (requires Enterprise)

See also: [\[\[ADR-007\]\]](#), [\[\[ADR-008\]\]](#), [\[\[ADR-018\]\]](#)

---

## ADR-10

### Status

Accepted. Implemented in `platform-search/src/main/java/com/cogniterra/platform/search/application/qdrant/QdrantVectorSearchProvider.java`.

### Context

Vector search requires a database optimized for high-dimensional similarity queries. General-purpose databases (PostgreSQL) can support vectors via extensions (pgvector) but are not optimized for vector-first workloads.

### Decision

Use **Qdrant** as the dedicated vector database. Qdrant is purpose-built for vector similarity search with:

- Native cosine similarity scoring
- Payload filtering (metadata alongside vectors)

- Collection management with configurable vector dimensions
- REST + gRPC APIs

The `QdrantVectorSearchProvider` communicates via REST API ( `/collections/{name}/points/search` ). Each vector point carries a payload with `chunkId` , `documentId` , `title` , `documentType` , `category` , `tags` , and `source` .

Qdrant is **optional**: when `platform.search.qdrant.host` is not configured, `NoOpVectorSearchProvider` provides empty search results and keyword search continues.

Configuration is managed via `QdrantProperties` ( `@ConfigurationProperties(prefix = "platform.search.qdrant")` ). The default collection is `document_intelligence_chunks` with 768-dimensional vectors (matching `nomic-embed-text` output).

## Alternatives Considered

- **pgvector (PostgreSQL extension)**: Rejected as primary vector store. While pgvector works for small-to-medium corpora, Qdrant provides better performance at scale, native quantization, and is purpose-built for vector search. pgvector is used for development/testing via H2 compatibility.
- **Weaviate, Pinecone, Milvus**: Rejected. Qdrant was chosen for its Rust performance, simple deployment (single binary), and REST API. The `VectorSearchProvider` SPI makes switching straightforward.
- **In-memory vector search**: Rejected. Does not persist across restarts.

## Consequences

- **High-performance vector search**: Cosine similarity with payload filtering
- **Collection auto-creation**: `QdrantCollectionManager.ensureCollectionExists()` on startup
- **Batch indexing**: `indexBatch()` for efficient bulk ingestion
- **Graceful degradation**: Platform works without Qdrant (keyword-only search)

## Trade-offs

- Additional infrastructure dependency in production
- REST API adds ~5-10ms latency vs gRPC
- No built-in quantization or disk-based indexing in the current configuration

## Future Evolution

- gRPC client for lower latency
- Qdrant quantization for memory efficiency with large corpora
- Multi-collection strategy per tenant or document type
- Hybrid search pushdown to Qdrant (if Qdrant adds text search)

See also: [\[\[ADR-008\]\]](#), [\[\[ADR-011\]\]](#)

---

## ADR-11

### Status

Accepted. Implemented in `platform-ai/src/main/java/com/cogniterra/platform/ai/application/DefaultRetrievalOrchestrator.java` .

### Context

Different query types benefit from different retrieval strategies. A factual lookup ("What was the revenue in Q2?") benefits from keyword search. An exploratory question ("What capabilities does the platform have?") benefits from hybrid search. An

index inspection query should not waste resources on vector search.

## Decision

Implement a **Retrieval Orchestrator** that selects the retrieval strategy based on classified query intent:

User Query → Intent Classification → Strategy Selection → Search Execution → Result

Strategy mapping (deterministic and explainable):

| QueryIntent                                             | SearchMode | Rationale                                            |
|---------------------------------------------------------|------------|------------------------------------------------------|
| INDEX_INSPECTION, CORPUS_DISCOVERY                      | KEYWORD    | Exact match queries; vector search unnecessary       |
| QUESTION_ANSWERING, WORKSPACE_ANALYSIS, SOURCE_ANALYSIS | HYBRID     | Complex queries benefit from keyword + vector fusion |
| DOCUMENT_RESEARCH, DOCUMENT_LOOKUP                      | SEMANTIC   | Conceptual queries benefit from vector similarity    |

The orchestrator produces `RetrievalOrchestrationResult` with full **explainability metadata**: intent, strategy, mode, prompt template, result counts, fusion method, reranking status, timing, and a step-by-step `traceLog`.

## Alternatives Considered

- **Always run all retrievers:** Rejected. Wastes resources on inappropriate strategies (e.g., vector search for exact ID lookups).
- **User-specified strategy:** Rejected. Users should not need to understand retrieval internals.
- **ML-based strategy selection:** Deferred. A trained classifier could outperform keyword-based intent classification, but keyword rules are deterministic, explainable, and sufficient for initial release.

## Consequences

- **Deterministic:** Same query always produces same strategy
- **Explainable:** Every retrieval decision is recorded in `traceLog`
- **Efficient:** Only appropriate retrievers are executed
- **Prompt-aware:** Retrieves the latest prompt template from `PromptRegistry`

## Trade-offs

- Keyword-based intent classification can misclassify queries
- Strategy mapping is static (no runtime adaptation based on result quality)
- Graph retrieval is not automatically selected by the orchestrator (requires explicit `HYBRID_GRAPH` mode)

## Future Evolution

- ML-based intent classification for higher accuracy
- Feedback loop: if results are poor, retry with expanded strategy
- Graph strategy auto-selection when enriched entities are detected in the query

See also: `[[ADR-005]]`, `[[ADR-008]]`, `[[ADR-012]]`

---

## ADR-12

# Status

Accepted. Implemented in `platform-ai/src/main/java/com/cognitera/platform/ai/model/RetrievalOrchestrationResult.java` and `platform-ai/src/main/java/com/cognitera/platform/ai/model/InferenceMetadata.java` .

# Context

AI systems must be auditable. When the platform generates an answer, stakeholders need to know: which model generated it, which prompt template was used, which documents were retrieved, which retrieval strategy was selected, and how long each step took. Without explainability, AI output is a black box.

# Decision

Make **explainability a first-class concern** on every inference. Every AI response carries `InferenceMetadata` and every retrieval carries `RetrievalOrchestrationResult` with:

| Metadata                                  | Example                                                      |
|-------------------------------------------|--------------------------------------------------------------|
| provider                                  | "ollama"                                                     |
| model                                     | "qwen2.5:14b"                                                |
| promptTemplateId                          | "rag-answer/v1"                                              |
| retrievalStrategy                         | "HYBRID"                                                     |
| keywordResultCount                        | 12                                                           |
| vectorResultCount                         | 8                                                            |
| graphNodeCount                            | 3                                                            |
| totalSourceCount                          | 15                                                           |
| fusionMethod                              | "weighted-linear-fusion"                                     |
| rerankingApplied                          | true                                                         |
| rerankingProvider                         | "ollama-cross-encoder"                                       |
| retrievalStartedAt / retrievalCompletedAt | timestamps                                                   |
| traceLog                                  | ["Orchestration started", "Intent: QUESTION_ANSWERING", ...] |
| evaluationScores                          | {grounding: 0.72, faithfulness: 0.85}                        |

The `explain()` method produces a human-readable summary:

```
Intent: QUESTION_ANSWERING | Strategy: HYBRID | Prompt: rag-answer/v1 | Sources: 15 | Fusion:
weighted-linear-fusion | Reranking: yes (ollama-cross-encoder) | Duration: 245ms
```

Explainability is **built into the orchestration layer**, not bolted on as an afterthought. Business services don't call explainability APIs — the orchestrator populates metadata automatically.

# Alternatives Considered

- **Optional explainability:** Rejected. Explainability should never be optional in an enterprise AI platform.

- **Separate explainability service:** Rejected. Would require duplicating orchestration state. Building metadata into the orchestration result ensures consistency.
- **User-facing explainability only:** Rejected. Internal diagnostics are as important as user-facing explanations.

## Consequences

- **Every inference is auditable:** Full traceability from query to answer
- **Deterministic:** Same query → same strategy → same explainability output
- **Low overhead:** Metadata is collected during normal execution, not as a separate pass
- **Future-proof:** New retrieval sources automatically appear in metadata

## Trade-offs

- `traceLog` is append-only string list (not structured log entries)
- `RetrievalOrchestrationResult` uses a builder with 19 setters (verbose but explicit)
- Metadata is not yet exposed through a dedicated API endpoint

## Future Evolution

- Structured trace events with typed metadata (not string list)
- Explainability REST API ( `GET /api/inferences/{id}/explain` )
- Visualization of retrieval decisions (Sankey diagram of query → strategy → results)
- Differential explainability: "why was document A ranked above document B?"

See also: [\[\[ADR-011\]\]](#), [\[\[ADR-013\]\]](#), [\[\[ADR-014\]\]](#)

# ADR-13

## Status

Accepted. Implemented in `platform-ai/src/main/java/com/cogniterra/platform/ai/application/DefaultEvaluationService.java` .

## Context

AI-generated answers must be evaluated for quality. Without evaluation, there is no feedback loop to detect hallucination, poor grounding, or irrelevant answers. Evaluation must run automatically as part of the inference pipeline, not as a separate manual process.

## Decision

Implement an **Evaluation Engine** that automatically evaluates every retrieval + inference cycle:

| Metric                               | Method                                    | Range       |
|--------------------------------------|-------------------------------------------|-------------|
| <code>groundingScore</code>          | Context length vs answer length ratio     | [0, 1]      |
| <code>citationCoverage</code>        | Citation markers [1], [2] in answer       | [0, 1]      |
| <code>faithfulness</code>            | Inverse of hallucination indicators       | [0, 1]      |
| <code>answerRelevance</code>         | Term overlap between question and answer  | [0, 1]      |
| <code>contextRelevance</code>        | Term overlap between question and context | [0, 1]      |
| <code>hallucinationIndicators</code> | Uncertainty phrase detection              | integer ≥ 0 |

|        |                        |         |
|--------|------------------------|---------|
| passed | Composite quality gate | boolean |
|--------|------------------------|---------|

The `DefaultEvaluationService` uses heuristic methods (regex patterns, term overlap, length ratios). In production, a dedicated evaluation LLM or framework (deepeval, ragas) would replace these heuristics.

Evaluation is wired into `DefaultRetrievalOrchestrator` — it runs after retrieval and before the result is returned. Every `RetrievalOrchestrationResult` carries evaluation scores in its metadata.

Alternatives Considered

- **No evaluation:** Rejected. An AI platform without quality feedback is irresponsible engineering.
- **LLM-as-judge only:** Deferred. Requires a second LLM call, adding latency and cost. Heuristic evaluation provides fast, deterministic feedback. LLM evaluation can be added as a separate evaluation profile.
- **Human evaluation only:** Rejected. Does not scale and cannot run in CI pipelines.

Consequences

- **Automatic quality gate:** Every retrieval is scored; low-quality retrievals are flagged
- **Deterministic:** Heuristic methods produce consistent, reproducible scores
- **Lightweight:** Evaluation adds negligible latency (no LLM calls)

Trade-offs

- Heuristic methods are less accurate than LLM-based evaluation
- `hallucinationIndicators` uses regex patterns that miss sophisticated hallucinations
- No benchmark dataset integration for regression testing

Future Evolution

- LLM-as-judge evaluation (separate profile, gated by configuration)
- RAG evaluation framework integration (deepeval, ragas)
- Evaluation benchmark dataset with ground truth annotations
- Automated regression testing: "does this prompt change improve or degrade evaluation scores?"

See also: [\[\[ADR-011\]\]](#), [\[\[ADR-012\]\]](#), [\[\[ADR-014\]\]](#)

ADR-14

Status

Accepted. Implemented in `platform-ai/src/main/java/com/cognitera/platform/ai/application/DefaultWorkflowEngine.java` .

Context

Enterprise AI applications involve multi-step processes: document upload → extraction → enrichment → analysis → review → completion. These processes have defined steps, transitions, and state. Hardcoding step logic in controllers couples workflow to HTTP concerns. A reusable workflow engine separates process definition from execution.

Decision

Implement a reusable **Workflow Engine** ( `WorkflowEngine` interface):

| Concept | Implementation |
|---------|----------------|
|---------|----------------|

|                    |                                                                     |
|--------------------|---------------------------------------------------------------------|
| WorkflowDefinition | Named workflow with ordered steps and transitions                   |
| WorkflowStep       | Step with id, name, description, handler type (manual/automated/ai) |
| WorkflowInstance   | Running instance with current step, status, context                 |

Two pre-registered workflows demonstrate the engine:

1. **document-intelligence** : SETUP → INGESTION → ANALYSIS → REVIEW → COMPLETE (5 steps)
2. **batch-ingestion** : INGEST → ENRICH → COMPLETE (3 steps)

The engine supports:

- `start(definitionId, context)` — creates a new instance at the initial step
- `advance(instanceId)` — moves to the next step; marks COMPLETED when no further steps
- `previous(instanceId)` — returns to the previous step
- `findInstance(instanceId)` — retrieves instance state
- `listActive()` — finds all active instances

The current implementation is in-memory (suitable for single-node). A database-backed store would be needed for multi-node deployments.

## Alternatives Considered

- **Hardcoded wizard in controller:** Rejected. The old workspace wizard was a hardcoded switch statement in `WorkspacePageController`. A reusable engine separates process logic from presentation.
- **Camunda/Flowable BPMN engine:** Rejected. Over-engineered for the current use case. The lightweight engine demonstrates the concept without framework lock-in.
- **Spring State Machine:** Rejected. Adds framework dependency for a relatively simple state transition model.

## Consequences

- **Reusable:** Any module can define and execute workflows
- **Simple API:** 5 methods, intuitive lifecycle
- **Pre-registered workflows:** Demonstrates real usage without configuration files

## Trade-offs

- In-memory storage means workflows reset on restart
- No persistence or fault tolerance for long-running workflows
- No parallel step execution or conditional branching
- Step handler types are declarative (manual/automated/ai) but not enforced

## Future Evolution

- Database-backed instance storage for production durability
- Conditional transitions based on step outcomes
- Timer-based transitions (escalation after timeout)
- Visual workflow definition (BPMN subset)
- Workflow event publishing for audit integration

See also: [\[\[ADR-001\]\]](#), [\[\[ADR-020\]\]](#)

---

## ADR-15

## Status

Accepted. Implemented in `platform-observability/src/main/java/com/cognitera/platform/observability/metrics/AiMetrics.java`.

## Context

AI operations are complex, multi-step processes with multiple external dependencies. Without observability, diagnosing failures (is the LLM slow? is Qdrant down? is enrichment producing too few entities?) requires log diving. Production AI systems require metrics.

## Decision

Implement **AI-specific observability** using Micrometer + Prometheus, integrated into a dedicated `platform-observability` module:

| Metric                                   | Type                         | Tags                           |
|------------------------------------------|------------------------------|--------------------------------|
| <code>ai.inference.duration</code>       | Timer (percentile histogram) | provider, model                |
| <code>ai.embedding.duration</code>       | Timer                        | provider, model                |
| <code>ai.retrieval.duration</code>       | Timer                        | mode (keyword/semantic/hybrid) |
| <code>ai.graph.retrieval.duration</code> | Timer                        | —                              |
| <code>ai.enrichment.duration</code>      | Timer                        | —                              |
| <code>ai.ingestion.duration</code>       | Timer                        | document_type                  |
| <code>ai.prompt.duration</code>          | Timer                        | template                       |
| <code>ai.provider.available</code>       | Gauge (0/1)                  | provider                       |
| <code>ai.embedding.count</code>          | Counter                      | provider                       |
| <code>ai.enrichment.entities</code>      | Counter                      | —                              |
| <code>ai.evaluation.*</code>             | Summary                      | metric name                    |

Metrics are exposed via `/actuator/prometheus` and `/actuator/health`. Health indicators check Ollama, Qdrant, and provider availability.

The `platform-observability` module is reusable across future applications. It depends only on Micrometer and Spring Boot Actuator — no platform-specific dependencies.

## Alternatives Considered

- **Logging only:** Rejected. Logs are not aggregatable or queryable at scale. Metrics enable dashboards and alerting.
- **OpenTelemetry from day one:** Deferred. Micrometer provides a simpler API that can export to OTLP when OpenTelemetry is adopted.
- **Metrics in each module:** Rejected. Centralizing metrics in `platform-observability` ensures consistent naming, avoids duplication, and enables reuse.

## Consequences

- **Operational visibility:** Every AI operation produces metrics

- **Reusable module:** Future applications get observability by depending on `platform-observability`
- **Standard format:** Prometheus exposition format enables Grafana dashboards
- **Health endpoints:** Liveness, readiness, and dependency health checks via Actuator

### Trade-offs

- Metrics are defined but not yet fully wired into all services (  `AiMetrics` is injected in fewer places than ideal)
- No custom Grafana dashboard definitions
- No alerting rules (would be added in deployment configuration, not in the platform)

### Future Evolution

- Wire  `AiMetrics` into all AI services (  `AiService` ,  `SearchService` ,  `DefaultEnrichmentService` )
- OpenTelemetry exporter for distributed tracing
- Pre-built Grafana dashboard JSON
- Alert definitions for critical metrics (LLM latency > threshold, provider down)

See also: `[[ADR-002]]`, `[[ADR-016]]`

## ADR-16

### Status

Accepted. Implemented throughout the platform via  `@ConditionalOnProperty` ,  `@ConditionalOnBean` ,  `ObjectProvider` , and  `NoOp*` fallback beans.

### Context

The Enterprise AI Platform depends on multiple external services: Ollama (LLM), Qdrant (vector DB), Neo4j (graph DB), PostgreSQL (relational DB). In production, any of these may be unavailable. The platform must continue functioning with reduced capabilities rather than failing entirely.

### Decision

Design every external dependency as **optional**. The platform uses multiple Spring mechanisms for graceful degradation:

| Mechanism                             | Example                                                                                                    | Behavior                                 |
|---------------------------------------|------------------------------------------------------------------------------------------------------------|------------------------------------------|
| <code> @ConditionalOnProperty</code>  | <code> OllamaChatProvider</code> requires <code> platform.ai.ollama.base-url</code>                        | Bean not created if config absent        |
| <code> @ConditionalOnBean</code>      | <code> Neo4jGraphSearchAdapter</code> requires <code> GraphEnrichmentService</code>                        | Bean not created if Neo4j unavailable    |
| <code> ObjectProvider&lt;T&gt;</code> | <code> DefaultDocumentIngestionProcessor</code> injects <code> ObjectProvider&lt;EnrichmentHook&gt;</code> | Null-safe <code> getIfAvailable()</code> |
| <code> NoOp*</code> fallbacks         | <code> NoOpVectorSearchProvider</code> , <code> NoOpGraphSearchProvider</code>                             | Return empty results, never throw        |
| try-catch wrappers                    | <code> Neo4jGraphSearchAdapter.search()</code>                                                             | Log warning, return empty list           |
| <code> isAvailable()</code> checks    | <code> GraphSearchProvider.isAvailable()</code> ,<br><code> ChatCompletionProvider.isAvailable()</code>    | Caller checks before invoking            |

|                            |                                              |                                     |
|----------------------------|----------------------------------------------|-------------------------------------|
| ProviderRouter<br>fallback | 4-tier selection skips unavailable providers | Throws only when ALL<br>unavailable |
|----------------------------|----------------------------------------------|-------------------------------------|

The degradation is **visible**: health indicators report status, metrics track failures, logs record the reason. The platform never silently degrades.

### Alternatives Considered

- **Mandatory all services**: Rejected. Would prevent development without full infrastructure.
- **Circuit breaker only**: Rejected. Circuit breakers (Resilience4j) are valuable for transient failures but don't address the "service never configured" case. Conditional beans handle both.
- **Feature flags**: Rejected. Adds complexity. Conditional bean activation is simpler and more idiomatic in Spring.

### Consequences

- **Platform starts with zero infrastructure** (except PostgreSQL for basic operation)
- **157 tests pass without any external services** (H2 in-memory, no Ollama/Qdrant/Neo4j)
- **Clear failure modes**: When a service is down, the reason is logged and surfaced in health endpoints
- **Gradual capability ramp**: Start with keyword search → add Qdrant for vector → add Neo4j for GraphRAG → add Ollama for LLM

### Trade-offs

- Some operations silently return empty results (graph search) rather than surfacing errors to users
- No retry logic for transient failures (deferred to Resilience4j integration)
- Health indicators currently read env vars directly rather than Spring config (known drift, documented)

### Future Evolution

- Resilience4j integration for retry, circuit breaker, rate limiter
- Health indicators refactored to inject `@ConfigurationProperties`
- Degradation event publishing to audit log
- User-facing capability indicators ("vector search unavailable")

See also: [\[\[ADR-003\]\]](#), [\[\[ADR-008\]\]](#), [\[\[ADR-009\]\]](#), [\[\[ADR-010\]\]](#), [\[\[ADR-015\]\]](#)

## ADR-17

### Status

Accepted. Interface defined in `platform-ai/src/main/java/com/cognitera/platform/ai/api/DomainConfiguration.java`. Default implementations embedded in existing services.

### Context

The platform must support multiple application domains (contract intelligence, financial analysis, regulatory compliance, technical documentation) from a single codebase. Domain-specific logic (concept definitions, analysis objectives, finding hierarchies, system instructions) must not be hardcoded in platform services. The platform core must remain domain-independent.

### Decision

Define a `DomainConfiguration` **SPI** that encapsulates domain-specific AI behavior:

| Method                                  | Returns                   | Purpose                                     |
|-----------------------------------------|---------------------------|---------------------------------------------|
| <code>domainId() / displayName()</code> | String                    | Identity                                    |
| <code>concepts()</code>                 | List<ConceptDefinition>   | Keywords + governing references per concept |
| <code>objectives()</code>               | List<ObjectiveDefinition> | Analysis objectives with keywords           |
| <code>findingRoleMapping()</code>       | Map<String, String>       | Reference → finding role mapping            |
| <code>findingRelationships()</code>     | List<FindingRelationship> | Relationships between findings              |
| <code>centralityWeights()</code>        | Map<String, Double>       | Centrality weights for references           |
| <code>peripheralReferences()</code>     | Set<String>               | References classified as peripheral         |
| <code>roleKeywords()</code>             | Map<String, List<String>> | Source role classification keywords         |
| <code>systemInstruction()</code>        | String                    | Domain-specific AI system instruction       |
| <code>answerStructureGuidance()</code>  | String                    | Domain-specific answer structure            |

Applications provide their `DomainConfiguration` as a Spring bean. The platform's existing services ( `DefaultConceptExtractionService` , `DefaultObjectiveAnalysisService` , etc.) currently embed default configurations. These defaults represent the original document intelligence domain but should be migrated to `DomainConfiguration` implementations.

## Alternatives Considered

- **Hardcode domain logic in each service:** Rejected. Couples the platform to a single domain and prevents reuse.
- **Database-driven configuration:** Deferred. Adds infrastructure dependency. SPI-based configuration is simpler and testable.
- **Remove all domain logic from platform:** Rejected. The platform must demonstrate real AI behavior. `DomainConfiguration` keeps domain logic while making it swappable.

## Consequences

- **Extension point exists:** New domains implement `DomainConfiguration` without touching platform code
- **Migration path:** Existing hardcoded domain rules in 8 services can migrate to `DomainConfiguration` implementations incrementally
- **Testability:** Domain configurations can be tested independently

## Trade-offs

- Currently a forward-looking SPI — the 8 existing services still embed their configurations rather than consuming `DomainConfiguration`
- No multi-domain routing (which `DomainConfiguration` to use for a given query)
- No domain discovery mechanism

## Future Evolution

- Migrate 8 services to consume `DomainConfiguration`
- Multi-domain routing based on query classification
- Domain configuration discovery ( `List<DomainConfiguration>` injection for multi-domain setups)
- External domain configuration via YAML/JSON files

See also: [[ADR-003]], [[ADR-007]]

---

# ADR-18

## Status

Accepted. Implemented in `platform-neo4j/src/main/java/com/cognitera/platform/neo4j/model/GraphNode.NodeProvenance` .

## Context

Knowledge graphs generated from AI extraction are only trustworthy if every node and edge can be traced back to its source. Without provenance, a graph node claiming "Acme Corporation is an ORGANIZATION" cannot be audited — was this extracted from a financial report or hallucinated by an LLM?

## Decision

Every graph node and relationship carries `NodeProvenance` :

| Field                             | Purpose                                    |
|-----------------------------------|--------------------------------------------|
| <code>sourceDocumentId</code>     | Which document was the source              |
| <code>chunkId</code>              | Which chunk within the document            |
| <code>chunkOffset</code>          | Position within the chunk                  |
| <code>extractionConfidence</code> | How confident was the extraction (0.0–1.0) |
| <code>extractionTimestamp</code>  | When was it extracted                      |
| <code>extractionModel</code>      | Which model performed the extraction       |
| <code>promptVersion</code>        | Which prompt template version was used     |
| <code>provider</code>             | Which AI provider performed the extraction |

Provenance is captured during enrichment ( `DefaultEnrichmentService` ), bridged through `EnrichmentHook` , and persisted to Neo4j by `GraphEnrichmentService` . The data flows: `EnrichmentContext` → `EnrichmentResult` → `GraphNode` (with `NodeProvenance` ) → Neo4j.

## Alternatives Considered

- **No provenance:** Rejected. An unauditable knowledge graph is useless for enterprise applications.
- **Separate provenance table in PostgreSQL:** Rejected. Graph data should carry its own provenance. A separate store creates consistency challenges.
- **Blockchain-based provenance:** Rejected. Over-engineered. Immutable audit log in PostgreSQL + provenance in Neo4j provides sufficient traceability.

## Consequences

- **Full traceability:** Every graph element can be traced to its source document, chunk, model, prompt, and provider
- **Confidence-aware:** Extraction confidence enables filtering low-confidence extractions
- **Reproducibility:** Re-extraction with different models/prompts can be compared by timestamp and version

## Trade-offs

- Provenance adds storage overhead to every node and relationship

- `extractionConfidence` is self-reported by the extraction method (not independently verified)
- Provenance is not yet leveraged for retrieval scoring (e.g., boost high-confidence extractions)

### Future Evolution

- Confidence-weighted retrieval: boost documents with high-confidence extractions
- Provenance visualization in the UI
- Automated re-extraction when prompt version changes
- Provenance chain: "entity A was extracted from chunk B of document C by model D using prompt E at time F"

See also: [\[\[ADR-007\]\]](#), [\[\[ADR-008\]\]](#), [\[\[ADR-009\]\]](#)

## ADR-19

### Status

Accepted. Implemented across `platform-ai/src/test/` and `platform-api/src/test/`. Documented in `TESTING.md`.

### Context

An Enterprise AI Platform requires confidence that every architectural subsystem behaves correctly. Tests must validate behavior, not implementation details. A reader should understand the platform architecture by reading the tests.

### Decision

Adopt a **5-layer testing strategy**:

| Layer           | Location                                             | Execution                          | Purpose                                 |
|-----------------|------------------------------------------------------|------------------------------------|-----------------------------------------|
| Unit            | <code>platform-ai/src/test/.../unit/</code>          | <code>mvn test</code>              | Validate components in isolation        |
| Integration     | <code>platform-api/src/test/.../ai/</code>           | <code>mvn verify</code>            | Validate subsystems with Spring context |
| Architecture    | <code>platform-api/src/test/.../architecture/</code> | <code>mvn verify</code>            | Validate end-to-end behavior            |
| Contract        | <code>platform-ai/src/test/.../contract/</code>      | <code>mvn verify</code>            | Validate SPI stability                  |
| Playwright (UI) | <code>e2e-tests/playwright/</code>                   | <code>mvn verify -Pui-tests</code> | Validate browser behavior               |

Key principles:

- **Behavior over implementation:** Tests verify what the platform does, not how it does it
- **Executable documentation:** Test names describe behavior ( `"routes openai:gpt-4o to openai provider"` )
- **Realistic scenarios:** Architecture tests use a `test-corpus/` with real document types and expected behaviors
- **Contract tests:** Every `ChatCompletionProvider` implementation must satisfy the abstract `ChatCompletionProviderContract`
- **No mocks for platform internals:** Unit tests mock external providers; integration tests use real Spring context with H2

157 tests validate: Prompt Registry (12), Model Capability Registry (12), Provider Router (14), Provider Contracts (5), Retrieval Orchestrator (3), Evaluation Engine (8), Workflow Engine (10), Semantic Enrichment (14), Graceful Degradation (9), GraphRAG

(4), AI Pipeline (10), Browser UI (12), Platform (91).

## Alternatives Considered

- **Coverage-driven testing:** Rejected. Coverage percentage does not measure architectural confidence. 157 behavioral tests provide more value than 500 getter tests.
- **BDD framework (Cucumber):** Rejected. Adds framework complexity. JUnit 5 `@DisplayName` + `@Nested` provides sufficient behavior description.
- **No UI tests:** Rejected. Browser tests validate user-visible behavior that unit tests cannot.

## Consequences

- **Architecture as tests:** Test structure mirrors platform architecture
- **Confidence, not coverage:** Every major subsystem has executable proof of correctness
- **CI-ready:** `mvn clean verify` runs all tests except Playwright (separate profile)
- **Contract protection:** New providers must satisfy SPI contracts

## Trade-offs

- Playwright tests require a running application (separate Maven profile)
- `test-corpus/` is small (3 documents) — sufficient for architectural validation, not exhaustive
- No performance regression tests yet

## Future Evolution

- Performance smoke test suite ( `mvn verify -Pperformance` )
- Extended test corpus with multilingual documents
- Snapshot tests for prompt rendering output
- Automated evaluation regression: "did this prompt change degrade faithfulness?"

See also: [[ADR-020]], [TESTING.md](#)

---

## ADR-20

### Status

Accepted. Embodied in every architectural decision across the platform.

### Context

The platform is not a chatbot. It is not a RAG demo. It is a reference implementation of an Enterprise AI Platform designed to be reusable across multiple AI-powered applications. Every architectural decision flows from this philosophy.

### Decision

The platform is governed by these design principles:

#### 1. AI is Infrastructure

AI is not a feature. It is infrastructure, like databases or message queues. Business logic depends on abstract interfaces ( `ChatCompletionProvider` , `EmbeddingProvider` ), not concrete implementations ( `OllamaChatProvider` , `OpenAiChatProvider` ). Adding a new AI provider requires zero changes to business logic.

#### 2. Modular Monolith Over Microservices

Module boundaries are enforced at compile time. Clear dependency direction prevents cycles. The architecture supports future extraction into microservices but does not prematurely distribute.

### 3. Graceful Degradation is Mandatory

Every external dependency is optional. The platform starts with zero infrastructure and gains capabilities as services become available. Keyword search works without Qdrant. Retrieval works without Neo4j. Inference works without Ollama.

### 4. Explainability by Default

Every AI operation produces auditable metadata: which model, which prompt, which strategy, which documents, how long it took. Explainability is not a feature toggle — it is built into the orchestration layer.

### 5. Documents Are the Knowledge Source

Knowledge is extracted from documents automatically during ingestion. There is no manual knowledge base, no CRUD interface for entities. The enrichment engine extracts entities, concepts, and relationships. The knowledge graph is auto-generated.

### 6. Domain Independence

The platform core contains no domain-specific logic. Domain customization happens through `DomainConfiguration` implementations. The platform can serve contract intelligence, financial analysis, compliance, or technical documentation from the same codebase.

### 7. Testing as Architecture Documentation

Tests validate behavior, not implementation. Test structure mirrors platform architecture. A Principal Engineer should understand the platform by reading the test suite.

### 8. Production Readiness

Every subsystem considers: how does it fail? How is it observed? How is it configured? Conditional beans, health indicators, metrics, structured logging, and configuration properties are not afterthoughts — they are first-class concerns.

## Alternatives Considered

The alternative philosophies considered and rejected:

- **"Move fast and break things"**: Incompatible with enterprise AI. Every decision is deliberate, documented, and testable.
- **"Maximize features, minimize architecture"**: Would produce a prototype, not a platform. The architecture is the product.
- **"AI is magic"**: Treating AI as opaque magic leads to unmaintainable systems. Every AI operation is instrumented, evaluated, and auditable.

## Consequences

- **Reusable platform**: Future applications (contract intelligence, financial analysis, compliance) can be built on the same foundation
- **Clear extension points**: `DomainConfiguration`, `ChatCompletionProvider`, `GraphSearchProvider` define where the platform grows
- **Professional engineering**: The codebase demonstrates architecture, testing, observability, and resilience — not just AI integration
- **9 modules, 157 tests, 0 failures**: Every architectural subsystem participates in real execution paths

## Trade-offs

- Architectural rigor means more interfaces and abstractions than a prototype would have
- The platform is not optimized for the shortest path to a demo — it is optimized for long-term maintainability
- Some SPLs ( `DomainConfiguration` ) are forward-looking and not yet consumed by all eligible services

## Future Evolution

This document should remain stable. New ADRs should reference these principles. If a proposed change conflicts with a principle, it should either be rejected or this ADR should be amended with explicit rationale.

See also: [\[\[ADR-001\]\]](#), [\[\[ADR-003\]\]](#), [\[\[ADR-012\]\]](#), [\[\[ADR-016\]\]](#), [\[\[ADR-017\]\]](#), [\[\[ADR-019\]\]](#)

---