

Enterprise AI Platform — Architecture & Engineering Handbook

Second Edition — June 2026

This handbook is for the engineer who looks at an AI demo and immediately asks: What happens when it fails? How do I extend it? Why was it designed this way?

Contents

Part I — Vision

1. [Enterprise AI Is Different](#)
2. [The Principles That Shape Every Decision](#)
3. [Why This Platform Exists](#)

Part II — Architecture

4. [How the Platform Fits Together](#) 5. [The Document Journey](#) 6. [How AI Is Wired In](#)

Part III — Subsystems

7. [Semantic Enrichment](#) 8. [GraphRAG](#) 9. [The Retrieval Decision Layer](#) 10. [Prompts as Engineering Assets](#) 11. [Models as Infrastructure](#) 12. [Workflows Without Hardcoding](#)

Part IV — Operations

13. [The Audit Trail](#) 14. [Observing What Happens](#) 15. [Security](#) 16. [How We Test This](#) 17. [Deployment](#)

Part V — Reflection

18. [Extending the Architecture](#) 19. [Performance Notes](#) 20. [The Trade-offs We Accepted](#) 21. [What We Learned](#) 22. [The Architecture Endures](#)

Appendices: [A — Key Decisions](#) · [B — Glossary](#) · [C — References](#)

Part I — Vision

1. Enterprise AI Is Different

There are two kinds of AI projects. The first kind — by far the more common — starts with a notebook. Someone loads a dataset, picks a model, writes a prompt, and gets impressive results within hours. The demo works. Everyone is excited.

Then an engineer asks a question the notebook never anticipated.

Can we trace which document this answer came from?

What happens when the vector database is down?

How do we switch from Ollama to OpenAI without rewriting the retrieval code?

Which version of the prompt produced this answer last Tuesday?

Can we prove to an auditor that the system didn't hallucinate this claim?

These are not unreasonable questions. They are exactly the questions enterprise software must answer. But they are invisible in a notebook. A notebook has one happy path. Enterprise software needs every path — especially the unhappy ones — handled deliberately.

This is the gap between an AI demo and an AI platform.

The Three Gaps

Three specific gaps separate demos from platforms.

The observability gap. A notebook produces output. A platform produces output *and* an audit trail. Which model ran? Which prompt version was used? Which documents were retrieved? How confident is the answer? Without this metadata, AI output is indistinguishable from guesswork.

The provider coupling gap. A notebook calls `openai.ChatCompletion.create()`. A platform calls `chatProvider.complete(prompt)`. The difference is indirection — and indirection is what makes provider swaps possible without rewriting business logic.

The degradation gap. A notebook crashes when an API is unreachable. A platform degrades gracefully — keyword search continues when vectors are unavailable, enrichment falls back to regex when the LLM is down, retrieval works without the knowledge graph. Every external dependency is optional by design.

Concern	Demo	Platform
Provider coupling	Hardcoded API calls	Abstract interfaces with conditional activation
Failure handling	Exceptions propagate to user	Graceful degradation with fallback implementations
Prompt management	Strings in source code	Versioned registry with audit trails
Answer quality	Subjective assessment	Automated evaluation: grounding, faithfulness, hallucination
Audit trail	None	Full explainability metadata on every inference
Extensibility	Rewrite code	Implement SPI + add configuration

What This Platform Demonstrates

This platform closes those three gaps. It is a working reference implementation — not a product, not a framework, not a SaaS service. It demonstrates how a team of experienced software engineers integrates AI into infrastructure: with clear module boundaries, swappable providers, versioned prompts, automated evaluation, and graceful degradation baked into every external dependency.

The architecture embodies a simple thesis: **AI is infrastructure, not a feature**. Treat LLMs, embedding models, vector databases, and knowledge graphs the way you treat PostgreSQL — abstract them behind interfaces, select them via configuration, swap them without touching business logic.

That thesis shapes every decision that follows.

2. The Principles

Eight principles keep hundreds of design decisions coherent. When a proposal conflicts with a principle, the proposal is rejected — or the principle is amended explicitly, with documented rationale. There is no silent drift.

AI Is Infrastructure

LLMs, embedding models, vector databases, knowledge graphs — these are infrastructure components. They belong behind interfaces. Business logic depends on abstractions. Implementations activate based on configuration.

A document enrichment service never calls a specific LLM provider. It calls an interface. Whether that interface is backed by Ollama, OpenAI, or a regex fallback is not a business concern. This pattern repeats everywhere — retrieval, embedding, reranking, evaluation.

Graceful Degradation, Not Hard Failures

In production, external services fail. Designing for this from day one means every external dependency is optional. The platform starts with only PostgreSQL. Add Qdrant: vector search activates. Add Neo4j: the graph populates. Add Ollama: LLM inference runs. Remove any: the platform continues, reduced capabilities logged clearly, never silently.

This was not added after the fact. Conditional bean activation, fallback implementations, and availability checks are built into every provider boundary.

Explainability by Default

If this platform generates an answer, it can explain itself — which intent was classified, which retrieval strategy was selected, which prompt template and version was used, which provider and model ran inference, how many chunks and graph nodes were retrieved, and what evaluation scores resulted. All generated as a byproduct of execution. No separate "explain this" pass. Zero cost in the happy path.

Documents Are the Knowledge Source

Users upload documents. The platform builds the knowledge. There is no manual knowledge base. The enrichment engine extracts entities, concepts, and relationships automatically during ingestion. The knowledge graph is auto-generated. An earlier version included a manual knowledge base — it was removed. The document corpus is the only source of truth.

Domain Independence

The core contains no assumptions about any industry or document type. Domain customization lives behind a pluggable interface. The included Document Intelligence application provides defaults. New domains add their own configuration — no core changes.

Testing as Architecture Documentation

Tests describe behavior, not implementation. A Staff Engineer should understand the platform by reading the test suite. Test names read like specifications.

Production Readiness

Observability, health checks, and configuration are first-class concerns designed alongside the features they observe — not afterthoughts added before release.

Modular Monolith

Module boundaries enforced at compile time. Dependency flows one way. Higher modules depend on lower modules — never the reverse.

3. Why This Platform Exists

Most AI projects begin the same way. Someone writes a script that calls an LLM API. They demo it on a laptop with one document. It works. Everyone is excited.

Then the engineering questions begin.

What happens when the vector database is down? Which prompt version generated that answer? How do we add a new provider alongside the existing one without rewriting the retrieval code? Can we trace why this answer cited that document?

These questions expose the gap between an AI demo and an AI platform. A demo needs one happy path. A platform needs every unhappy path handled deliberately.

This project closes that gap. It demonstrates how experienced engineers integrate AI into enterprise infrastructure — not as a feature bolted onto an existing system, but as infrastructure designed from first principles.

The result is a modular monolith: 9 Maven modules, a single deployable, compile-time boundaries enforcing dependency direction. It is not a product. It is not a framework. It is a foundation — code another engineering team could adopt as the starting point for their own AI application.

What's in scope: multi-provider AI orchestration, hybrid search across keyword, vector, and graph sources, automatic semantic enrichment during ingestion, GraphRAG with knowledge graph traversal, versioned prompt management, a queryable model capability database, automated evaluation of every answer, full explainability metadata, a reusable workflow engine, production metrics, and 157 automated tests.

What's deliberately absent: microservices, multi-agent systems, model fine-tuning, user management, billing. These are valid concerns that would distract from the core architecture.

Part II — Architecture

4. How the Platform Fits Together




Figure 4.1. Nine modules in three layers. Dependency flows downward. The application module depends on everything. Nothing depends on the application module.

The Modular Monolith

Three layers, nine modules, one direction of dependency.

At the bottom sits the audit module — the only leaf. Every other module depends on it for immutable audit logging with correlation IDs that thread through every request.

The middle layer holds seven platform modules. Each has one responsibility. Authentication. Document lifecycle. Hybrid search. AI orchestration. Knowledge graph persistence. Workspace management. Observability. When a module grows large, its internal package structure provides further separation without adding Maven modules. The goal is clarity, not a specific module count.

At the top, the assembly module wires everything together — REST controllers, UI templates, DTOs, and the configuration that conditionally activates provider beans based on which infrastructure is available. Future applications would create their own assembly modules on the same platform foundation.

Approach	Benefits	Costs
Modular Monolith	Compile-time boundaries; single deployable; clear dependency direction	No runtime isolation; modules share JVM
Microservices	Independent deployment; isolated failure domains; per-service scaling	Network complexity; eventual consistency; operational overhead
Single JAR (no modules)	Simplest to build	No dependency enforcement; rapid architectural decay

External Dependencies




Figure 4.2. Four external systems. PostgreSQL is required. Everything else — Qdrant, Neo4j, LLM providers — is optional.




Figure 4.3. A single Spring Boot process communicates with its dependencies over standard protocols. Optional services have dashed borders — the architecture expects them to sometimes be absent.

The platform communicates with four external systems. PostgreSQL is the only hard requirement. Qdrant provides vector search. Neo4j stores the knowledge graph. Ollama or OpenAI provide LLM inference and embeddings. Each additional service adds capability but is never required.

Spring Boot 3.3 with Java 21 was chosen for ecosystem breadth and enterprise familiarity — not because it is the fastest or most lightweight framework, but because a reference implementation should prioritize clarity over novelty.

5. The Document Journey

A document's journey through the platform touches text extraction, semantic enrichment, chunking, embedding generation, and three separate persistence stores — all while the user who uploaded it has moved on to other things.

The Asynchronous Choice

The key architectural decision: run ingestion asynchronously. Document processing takes seconds to minutes. HTTP request threads should not block that long. Upload creates a pending job and returns immediately. A scheduled worker picks it up.



Figure 5.1. The ingestion pipeline. Upload creates a pending job. A scheduled worker polls every 10 seconds. Extraction, enrichment, chunking, and embedding run sequentially. Results land in PostgreSQL, Qdrant, and Neo4j.

The pipeline is sequential by design, not by limitation. Parallel enrichment and embedding would improve throughput but make the code harder to trace. The reference implementation prioritizes clarity. Production deployments may add parallelism — the architecture does not prevent it.

Enrichment Before Chunking

One design choice that emerged during testing: enrichment runs before chunking, not after. Full document context produces better entity disambiguation than individual chunks. A name like "Apple" is ambiguous in a 200-word chunk but clear in a 5,000-word financial report. This is the kind of decision that looks obvious in retrospect but required testing to discover.

Degradation in Practice

Each step in the pipeline degrades independently. No embedding provider configured? Chunks are stored keyword-only, and search still works. No LLM available? Regex patterns extract entities with reasonable accuracy. No Neo4j? Enrichment still runs — only graph persistence is skipped. The pipeline is designed as a chain of independent, fallible steps, not a monolithic transaction.

6. How AI Is Wired In

Answering a question with AI is not one operation. It is twelve — and each produces metadata that feeds the explainability trail.




Figure 6.1. The twelve-step inference pipeline. Every step produces metadata. Validation, grounding, and evaluation run even when earlier steps produce warnings — the platform surfaces concerns rather than hiding them.

The Full Pipeline

Intent classification. Before launching expensive retrieval, the system classifies the query. An index inspection routes to keyword-only. A substantive question enters the full RAG pipeline. This classification is deterministic — same query, same path, every time.

Strategy selection. The classified intent maps to a retrieval strategy. Keyword for exact lookups. Hybrid for complex questions. Semantic for conceptual exploration. The mapping is a simple lookup table, not a machine learning model — explainable, auditable, and sufficient for current query types.

Retrieval. Search fans out to keyword, vector, and optionally graph sources. Results are fused through weighted linear combination and reranked.

Context assembly and prompt building. Retrieved chunks and analysis objectives are assembled. The prompt registry resolves the latest template. Variables are substituted.

Provider selection. The router selects which LLM backend to use based on model prefix, capability requirements, preference, and availability.

LLM inference, validation, grounding. The selected provider generates raw answer text. Temporal consistency is checked. Claims are validated against available sources. Sources are re-attributed with confidence scores.

Evaluation and explainability. Every decision from the pipeline is captured — traceable, auditable, deterministic. The `explain()` method produces a human-readable summary.

The Provider Architecture




Figure 6.2. Five service provider interfaces decouple business logic from infrastructure. Implementations activate conditionally. Fallbacks ensure the platform works without any optional service.

Approach	Example	Consequence
Provider SPI	<code>chatProvider.complete(prompt)</code>	Swap providers without changing orchestration
Direct API calls	<code>openaiClient.chat(prompt)</code>	Every provider change requires rewriting business logic
Factory pattern	<code>ProviderFactory.get("openai")</code>	Centralized but still couples factory to provider names

The provider architecture is the Strategy pattern applied to AI infrastructure. Business logic depends on interfaces. Implementations activate based on configuration. The two never meet in source code. This is not novel — but applying it consistently to AI, where most projects hardcode provider specifics, is what makes the architecture worth studying.

Part III — Subsystems

7. Semantic Enrichment

Keyword search finds exact matches. Vector search finds semantically similar passages. Neither understands that "Acme Corporation" is an organization, or that "\$45.2 million" relates to "revenue," or that "Dr. Sarah Chen" signed a document on a specific date.

This structured understanding — who, what, when, how much — is what enables intelligent retrieval. Without it, a query like "What contracts did Acme sign in Q2?" relies on the hope that relevant documents contain those exact words near each other.

Approach	How it works	Limitation
Keyword search	Exact token matching	Misses synonyms, entities, relationships
Vector search	Semantic embedding similarity	Misses structured relationships (entities, concepts)
Semantic enrichment	Extract entities, concepts, relationships	Requires extraction accuracy; builds structured knowledge

 Semantic Enrichment Pipeline

Figure 7.1. Enrichment runs after text extraction, before chunking. Full document context improves entity disambiguation. Results flow to the knowledge graph when available.

Dual-Mode Extraction

The enrichment engine uses two modes, selected automatically: LLM-based extraction when a chat provider is available (higher accuracy, understands context), and regex pattern matching when no LLM is available (catches organizations, persons, dates, monetary amounts — enough for basic retrieval enhancement).

Mode	When used	Accuracy	Dependency
LLM-based	Chat provider available	Higher — understands context	Requires LLM
Regex	No LLM available	Baseline — patterns only	No dependency

Why Enrich During Ingestion

Enrichment is expensive — LLM calls, regex passes over full documents. Query-time latency budgets are measured in milliseconds. Pre-computing enrichment during ingestion means retrieval is fast and enrichment quality is known before any question is asked. This is an architectural trade-off: ingestion latency for query-time speed.

Provenance

Design Rationale: Every graph node carries provenance. Without it, "Acme Corporation is an ORGANIZATION" could be a hallucination. With provenance, you trace it to the exact document, chunk, model, prompt version, and timestamp. The knowledge graph is auditable — not in theory, in every node, by construction.

8. GraphRAG

Traditional RAG retrieves chunks by keyword and vector similarity. It works when relevant documents share vocabulary or embedding proximity with the query. It fails when documents are semantically related but textually distant.

Example: a contract signed by Acme Corporation and a financial report mentioning Acme's Q2 revenue share no keywords, and their embeddings may sit far apart in vector space. But they are undeniably related — through the entity "Acme Corporation."

Approach	Strength	Weakness
Keyword retrieval	Exact matches; fast	Misses synonyms, conceptual connections
Vector retrieval	Semantic similarity	Misses structured entity relationships
Graph traversal	Entity/concept connections	Requires populated knowledge graph

A knowledge graph captures these relationships explicitly. Traversing it reveals connections that keyword and vector search miss. This was the motivation for GraphRAG — not theoretical elegance, but a real retrieval blind spot.




Figure 8.1. Three retrieval sources converge in weighted fusion. Graph results boost existing candidates — they don't compete with them. When the graph is unavailable, retrieval degrades to keyword + vector.

The Central Design Choice

Graph results **boost, not replace**. A graph traversal finding that adds up to 15% score increase for candidates with related graph nodes. A document at position #15 by keyword + vector might move to #5 because it mentions the same entities as a top-ranked document. Graph context nudges rankings rather than overriding them.

A Cautionary Tale

An honest admission: GraphRAG was documented as a feature for months before it was implemented. The architecture manual described knowledge graph participation in retrieval. The code used a permanent no-op fallback. Building the adapter that bridged enrichment output to the retrieval SPI closed that gap. The lesson: document what *is* built, not what *should be* built. Documentation ahead of implementation creates confusion. Documentation trailing implementation is always accurate.

9. The Retrieval Decision Layer

Running every retriever for every query wastes resources. An index inspection does not need vector search. A conceptual question benefits from semantic similarity. A complex analytical question benefits from hybrid search with graph augmentation.

The solution: a decision layer whose sole responsibility is choosing which retrievers to use based on classified intent.




Figure 9.1. Intent drives strategy deterministically. Every routing decision is logged for auditability.

Query Type	Strategy	Why
Index inspection, corpus discovery	Keyword only	Exact match queries; vector search adds no value
Question answering, analysis	Hybrid	Complex queries benefit from keyword + vector fusion
Document research, lookup	Semantic	Conceptual queries match better on meaning than keywords

Deterministic Over ML-Based

The intent classifier uses keyword rules, not machine learning. This is deliberate. Deterministic routing means the same query always takes the same path — debuggable, auditable, predictable. A trained classifier would be more accurate but less transparent.

The consequence: misclassification is possible and propagates deterministically. A query classified as keyword-only when it needs hybrid search will produce poorer results every time. This is an explicit trade-off. If classification accuracy becomes critical, the classifier can be replaced without touching the orchestrator — the architecture supports it.

The Fusion Algorithm

Weighted linear combination — not Reciprocal Rank Fusion. RRF requires rank positions from every retriever, which are not always available (graph search returns discovery results, not ranked lists). Weighted linear fusion works with raw scores and is simpler to explain. The algorithm is accurately named — we don't claim to implement RRF.

```
score = (keyword×0.40 + vector×0.40 + confidence×0.20) × docTypeWeight + graphBoost
```

10. Prompts as Engineering Assets

Approach	Versioning	Audit trail	Discoverability	A/B testing
Inline strings in code	Git history only	Manual	Grep	Impossible
External files	File-level	Manual	Directory listing	Manual
Versioned registry	Per-prompt	Automatic	Query by category	Built-in

Prompts scattered across source code as string constants are invisible, unversioned, and untestable. When a prompt changes, there is no audit trail. When a new domain requires different instructions, there is no mechanism to provide them without changing code.

The prompt registry treats prompts as engineering assets — versioned, categorized, discoverable, auditable. Every prompt has a qualified ID recorded in inference metadata. Nine categories organize prompts by purpose. The `render()` method substitutes variables — simple, predictable, no template engine dependency.

The registry is deliberately simple. No template inheritance. No dependency between prompts. No runtime compilation. Complexity in prompt management should be justified by need, not added preemptively. The current design handles versioning, categorization, and rendering. When those prove insufficient, the interface can evolve — but the extension point is clear and the migration path is short.

11. Models as Infrastructure

Different models have different capabilities. GPT-4o supports vision and structured output. nomic-embed-text only does embeddings. qwen2.5:7b is fast and local but lacks vision. Business logic must not hardcode assumptions about model capabilities.

The model capability registry describes every known model across eight capability dimensions plus quantitative metadata: context window, output tokens, latency estimates, cost estimates, temperature, use cases.

Selection Strategy	How it works	When it applies
Model prefix	<code>openai:gpt-4o</code> routes to openai provider	Explicit model choice
Capability match	Request STREAMING → find capable provider	Capability requirement
Preferred provider	Caller preference	Known provider preference
First available	First healthy provider	Fallback

The router receives all available implementations via dependency injection. Adding a provider means implementing the interface and adding a Spring component — the router discovers it without code changes. This is the Strategy pattern applied to AI infrastructure.

The architecture prevents business logic from selecting providers directly. Services request capabilities — streaming, embeddings, JSON output — and the router maps capabilities to providers. This indirection is what makes adding providers a zero-change operation for orchestration code.

12. Workflows Without Hardcoding

Document intelligence involves multi-step processes. Upload leads to extraction, enrichment, chunking, indexing, review, completion. Hardcoding these steps in a controller couples process logic to HTTP concerns and prevents reuse.

Approach	Flexibility	Complexity	When to use
Controller logic	Low — hardcoded	Low	Single fixed workflow
Workflow engine	High — configurable	Medium	Multiple workflows, reusable
BPMN engine	Very high	High	Complex branching, enterprise BPM

The workflow engine separates process definition from execution. Five methods cover the API. Two workflows come pre-registered. Steps declare handler types as extension points.

The engine deliberately avoids BPMN frameworks or state machine libraries. Five methods on an interface. Deterministic linear transitions. Not the most powerful workflow engine — the simplest one that demonstrates the concept. A production system might adopt Camunda or Temporal, and the interface was designed so that substitution would not require changing callers.

Part IV — Operations

13. The Audit Trail

If this platform generates an answer, it can explain itself. Not because someone added a feature — because the orchestration layer records every decision as it happens. Intent classification. Strategy selection. Prompt template and version. Provider and model. Result counts from every retriever. Fusion method. Reranking status. Evaluation scores.

The metadata is generated during normal execution. There is no separate "explain this" pass, no post-hoc analysis, no second LLM call. Every code path produces metadata by construction.

The design principle: explainability must have zero marginal cost. If explaining an answer requires additional computation, it will be disabled under load. By making explanation a byproduct of execution, it stays on by default, in every environment, forever.

14. Observing What Happens

A platform running AI workloads needs metrics designed for AI workloads — not just request counts and error rates, but inference latencies broken down by provider and model, embedding batch durations, retrieval timings by strategy, enrichment entity counts, evaluation score distributions.

The observability module provides Micrometer metrics for every AI operation. Timers with percentile histograms. Gauges for provider availability. Summaries for evaluation scores. All exported to Prometheus and surfaced through health checks.

The metric definitions stabilize first. Full integration follows. This ordering is deliberate: define the interface, validate with a few integrations, then roll out broadly. Wiring everything at once produces metrics nobody asked for.

15. Security

Authentication uses JWT with BCrypt-12 password hashing. Refresh tokens are hashed before storage — never plaintext — and rotated on use. Both stateless API authentication and session-based form login are supported.

Provider API keys are injected via environment variable substitution — never hardcoded, never committed. User input enters through variable placeholders in vetted prompt templates. Controller responses are HTML-escaped. The security model is standard, proven, and deliberately unremarkable. Innovation in security is a bug, not a feature.

16. How We Test This



Testing Pyramid

Figure 16.1. Five testing layers. Default build runs everything except browser tests, which use a separate profile to keep the primary build fast.

Layer	Count	Validates
Unit	44	Individual components — registries, router, evaluation
Integration	91	Subsystem interaction with real Spring context
Architecture	10	End-to-end behavior — enrichment, retrieval, evaluation
Resilience	9	Graceful degradation — fallback, unavailable services
Contract	5	SPI stability — every provider satisfies abstract contract
Browser	12	User journeys via Playwright (separate profile)

The test suite is structured as executable architecture documentation. Tests describe behavior: "routes openai:gpt-4o to openai provider" — not "testProviderRouter." A Staff Engineer should understand the platform by reading the tests.

17. Deployment

The platform starts with `docker compose up -d` for PostgreSQL and Qdrant. Add `--profile graph` for the knowledge graph database. One command starts the application.

Configuration uses a single file with environment variable substitution. Every custom property uses the `platform.*` namespace. The deployment model is deliberately simple: one process, one configuration file, one command. A reference implementation should be easy to run on a development machine.

 Deployment Architecture

Figure 17.1. Production deployment. The knowledge graph database is in a separate Docker profile. Every additional service adds capability but is never required.

Part V — Reflection

18. Extending the Architecture

The architecture was designed for extension. Four patterns cover the most common needs.

Adding a provider. Implement the chat provider interface. Add component annotation with conditional property activation. Register model capabilities. The router discovers it automatically.

Adding a domain. Implement the domain configuration interface. Provide it as a Spring bean. Concepts, objectives, hierarchies, and instructions become available without core changes.

Adding a workflow. Define steps and transitions. Call the workflow engine's start method. The engine handles state.

Adding a prompt. Register a template. It is immediately versioned, categorized, and discoverable.

The goal was never to predict every extension point. It was to establish patterns — interface, conditional activation, registry, dependency injection — that make extension feel natural rather than like working around constraints.

19. Performance Notes

This is a reference implementation, not a production deployment. It has not been benchmarked. What follows are architectural observations.

Chunking scales linearly with document length. Embedding generation is sequential — parallelizable via virtual threads. Weighted fusion is linear in candidate count. Graph traversal is bounded at 2 hops. LLM inference timeout is configurable.

No caching layer exists. Each retrieval re-executes. This is appropriate for demonstrating architecture. It would be the first optimization in production.

20. The Trade-offs We Accepted

These decisions reflect engineering judgment, not universal truth. Each was debated, documented, and deliberately accepted.

Decision	Chosen	Alternative	Why
Architecture	Modular monolith	Microservices	Compile-time boundaries sufficient; extraction possible
Framework	Spring Boot 3.3	Quarkus	Ecosystem breadth; enterprise familiarity
Graph database	Neo4j property graph	RDF / relational	Traversal queries; schema flexibility
Fusion algorithm	Weighted linear	Reciprocal Rank Fusion	Heterogeneous scores; simpler explanation
Prompt management	Versioned registry	Inline strings	Audit trails; regression testing
Enrichment	Dual-mode (LLM + regex)	LLM only	Works without AI infrastructure
Provider routing	Deterministic 4-tier	ML-based	Explainable; auditable
UI testing	Playwright	Selenium	Modern API; auto-waits; reliability

Each decision has a dedicated ADR in the companion volume with deeper rationale, alternatives considered, and future evolution paths.

21. What We Learned

Clean interfaces make domain-specific code easy to remove. The original codebase contained 55 German BGB paragraph references and 25+ legal keywords embedded in enrichment and analysis services. They were removed without architectural changes because the interfaces had clean separation from domain logic. The lesson generalizes: invest in interface design early. It pays off when requirements shift. This is not theoretical — it happened, and the interfaces saved months of refactoring.

Unused dependencies accumulate invisibly. Five JARs from a Spring AI integration sat on the classpath for months with zero imports anywhere in the codebase. A monthly dependency analysis run would catch this — it costs nothing and prevents deadweight from accumulating. The build audit that discovered this also found Flyway dependencies (migrations written but never enabled) and transitive JARs pulled in by unused direct dependencies.

Bean wiring must be verified, not assumed. The entire AI inference pipeline — twelve steps, a dozen collaborating services — compiled cleanly and passed every test. But the pipeline was never connected to the controller. The AI page showed retrieval results without generating LLM answers. A simple wiring verification test caught this. In Spring applications, test that your beans are actually injected and called — not just that they compile.

Documentation must trail implementation, not lead it. GraphRAG appeared in the architecture manual months before the code implemented it. The knowledge graph was populated but never queried during retrieval. A permanent no-op fallback sat where real graph traversal should have been. Building the adapter closed that gap. Now documentation reflects code, not aspiration.

Prompt versioning is table stakes for any AI platform. What started as a simple registry — store templates, substitute variables — became the foundation for deterministic behavior validation, A/B testing prompts, audit trails, and regression testing. Every inference records exactly which prompt version produced it. This was not planned from the start. It became obvious the first time a prompt change caused a behavior regression and nobody could trace which answers used the old version.

Tests as documentation works. Evolving from wiring-verification tests to behavior-driven tests made the suite readable by architects who never touched the implementation. "Routes openai:gpt-4o to openai provider" communicates more than "testProviderRouter" — not because the test implementation differs, but because the name describes what the system does rather than what class it tests.

22. The Architecture Endures

Technologies change. Models evolve. Providers come and go. Vector databases are replaced by faster alternatives. LLM APIs shift with each release. Prompt engineering is still nascent. The specific choices made today — Ollama over vLLM, Qdrant over Weaviate, Neo4j over RDF — will look dated within a few years.

But the architectural principles that shaped those choices do not date.

Treating AI as infrastructure rather than a feature means the interface between business logic and AI backends remains stable even as the backends themselves change. Building for graceful degradation means the platform is never more fragile than its most unreliable dependency. Explaining every decision means the system is auditable by construction, not by afterthought. Extracting knowledge from documents during ingestion rather than query time means enrichment quality can be verified before any question is asked.

These principles are independent of any specific technology. They apply whether the platform runs on Ollama or OpenAI, whether vectors live in Qdrant or pgvector, whether the graph lives in Neo4j or a property graph embedded in the application.

The architecture documented in this handbook is not the final answer. It is a snapshot of engineering judgment applied to a specific set of technologies at a specific point in time. But the principles — abstraction, degradation, explainability, provenance, domain independence — will remain useful long after the specific implementations have been replaced.

That is what makes this an architecture worth studying, not just a codebase worth reading.

Appendix A — Key Decisions

Twenty Architecture Decision Records document every major engineering decision. The complete volume — with introduction, decision index, architecture timeline, and cross-references — is available as [Architecture-Decision-Records.pdf](#).

ADR	Title	Status	Category
001	Modular Monolith	Accepted	Architecture
002	Spring Boot 3.3	Accepted	Technology
003	Provider Abstraction via SPI	Accepted	AI Infrastructure
004	Provider Router	Accepted	AI Infrastructure
005	Prompt Registry	Accepted	AI Infrastructure
006	Model Capability Registry	Accepted	AI Infrastructure
007	Semantic Enrichment Engine	Accepted	Knowledge
008	GraphRAG	Accepted	Knowledge
009	Neo4j as Knowledge Graph Persistence	Accepted	Knowledge
010	Qdrant as Vector Database	Accepted	Knowledge
011	Retrieval Orchestration	Accepted	Knowledge
012	Explainability by Default	Accepted	Quality
013	Evaluation Engine	Accepted	Quality
014	Workflow Engine	Accepted	Extensibility
015	AI Observability	Accepted	Operations
016	Graceful Degradation	Accepted	Operations
017	DomainConfiguration	Accepted	Extensibility
018	Provenance-Aware Knowledge Graph	Accepted	Knowledge
019	Multi-Layer Testing Strategy	Accepted	Quality
020	Design Philosophy	Accepted	Governance

Appendix B — Glossary

Capability — A feature an AI model supports: streaming, vision, JSON mode, embeddings, tool calling.

Chat Provider — The SPI for LLM backends. Every provider implements this interface.

Chunk — A segment of document text (~1200 characters) with sentence-boundary awareness and overlap, indexed for retrieval.

Citation Coverage — Fraction of AI answer claims that cite a retrieved source.

Context Window — Maximum tokens a model can process in a single request.

Domain Configuration — Pluggable interface for domain-specific rules: concepts, objectives, hierarchies, instructions.

Embedding — A 768-dimensional vector representation of text, generated by an embedding model.

Enrichment — Automatic extraction of entities, concepts, and relationships from documents during ingestion.

Evaluation — Automated quality assessment: grounding, faithfulness, hallucination indicators, relevance.

Explainability — Metadata describing how an inference was produced, generated as a byproduct of execution.

Faithfulness — How factually accurate an answer is relative to its source documents.

Fusion — Combining retrieval results from multiple sources into a single ranked list.

GraphRAG — RAG enhanced with knowledge graph traversal. Graph results boost, not replace.

Grounding — The degree to which an answer is supported by retrieved evidence.

Hallucination — AI-generated text not supported by source documents.

Hybrid Search — Retrieval combining keyword, vector, and optionally graph search.

Inference — The process of an LLM generating text from a prompt.

Intent — The classified purpose of a query: question answering, document lookup, index inspection.

Knowledge Graph — A network of entities, concepts, and relationships extracted from documents.

Model Capability Registry — Searchable database of model features, limits, and performance estimates.

Modular Monolith — Single deployable with compile-time module boundaries.

Provenance — Metadata linking data to its origin: source document, extraction method, model, prompt version, timestamp.

Prompt Registry — Versioned store of prompt templates organized by category.

Provider — An AI backend implementing a platform SPI.

Provider Router — Component that selects which provider to use. 4-tier deterministic selection.

RAG — Retrieval-Augmented Generation. Grounding LLM responses in retrieved documents.

Registry — A queryable, versioned store of known entities (prompts, models, capabilities).

Reranking — Reordering candidates: score sort followed by LLM cross-encoder if available.

Retrieval Strategy — Which retrievers to use, selected based on classified query intent.

SPI — Service Provider Interface. A pluggable contract for extending the platform.

Structured Output — LLM responses constrained to a format like JSON.

Token — A unit of text processed by an LLM (~0.75 words in English).

Vector — A numerical representation of text meaning, 768 dimensions in this platform.

Vector Database — Database optimized for high-dimensional similarity queries.

Workflow — A configurable multi-step process with defined state transitions.

Workflow Engine — Reusable component for executing workflows.

Appendix C — References

Architecture & Engineering

- **Spring Boot 3.3 Reference** — The platform's runtime framework
- **Martin Kleppmann, *Designing Data-Intensive Applications*** — Foundational thinking on data systems
- **Martin Fowler, *Patterns of Enterprise Application Architecture*** — Patterns for modular design
- **Simon Brown, *C4 Model*** — Architecture visualization approach

AI & Retrieval

- **Lewis et al. (2020), *RAG for Knowledge-Intensive NLP Tasks*** — Foundational RAG paper
- **Microsoft Research (2024), *GraphRAG: From Local to Global*** — Graph-enhanced retrieval
- **Neo4j Graph Data Science** — Graph algorithms for centrality and community detection
- **Qdrant Documentation** — Vector search with payload filtering

Operations

- **Micrometer** — JVM metrics instrumentation
- **Prometheus** — Metrics collection and alerting
- **Playwright** — Cross-browser automation

Companion Volumes

- [Architecture Decision Records](#) — 20 ADRs
- [Developer Guide](#) — Build, run, extend, test, contribute