

Enterprise AI Platform — Architecture & Engineering Handbook

Third Edition — June 2026

This handbook is for the engineer who looks at an AI demo and immediately asks: What happens when it fails? How do I extend it? Why was it designed this way? If you are that engineer, start here. By the end, you should understand not only how this platform works, but how to think about designing one yourself.

Contents

Part I — Vision

1. [Enterprise AI Is Different](#)
2. [The Principles That Shape Every Decision](#)
3. [Why This Platform Exists](#)

Part II — Architecture

4. [How the Platform Fits Together](#) 5. [The Document Journey](#) 6. [How AI Is Wired In](#)

Part III — Subsystems

7. [Semantic Enrichment](#) 8. [GraphRAG](#) 9. [The Retrieval Decision Layer](#) 10. [Prompts as Engineering Assets](#) 11. [Models as Infrastructure](#) 12. [Workflows Without Hardcoding](#)

Part IV — Operations

13. [The Audit Trail](#) 14. [Observing What Happens](#) 15. [Security](#) 16. [How We Test This Deployment](#) 17.

Part V — Reflection

18. [Extending the Architecture](#) 19. [Design Rules](#) 20. [Performance Notes](#) 21. [The Trade-offs We Accepted](#) 22. [What We Learned](#) 23. [The Architecture Endures](#)

Appendices: [A — Key Decisions](#) · [B — Glossary](#) · [C — References](#)

About This Project

This is not an AI framework. It is not LangChain, Spring AI, Haystack, or any of the libraries that make calling LLMs easier. Those tools solve a different problem: they provide APIs for building AI applications quickly.

This project demonstrates the architectural patterns that separate an AI demo from an AI platform.

The distinction matters. A framework makes it easy to call a model. A platform makes it possible to swap models without rewriting code, to trace every answer to its source, to prove to an auditor that the system did not hallucinate, to keep running when infrastructure fails. These are architecture concerns.

When we started building, we made an explicit choice: design for the unhappy path first. What happens when the vector database is down? Which prompt version produced that answer last Tuesday? How do we add a new provider without touching business logic? These questions shaped every decision that follows.

The result is a modular monolith — 9 Maven modules, a single deployable, compile-time boundaries enforcing dependency direction. It runs on Java 21 and Spring Boot 3.3. It is tested with 157 automated tests and documented with 20 Architecture Decision Records.

It is not a production SaaS service. It is not a chatbot. It is not a benchmark. It is a reference implementation — the kind of codebase another engineering team could adopt as the starting point for their own AI application.

Part I — Vision

1. Enterprise AI Is Different

There are two kinds of AI projects. The first kind — by far the more common — starts with a notebook. Someone loads a dataset, picks a model, writes a prompt, and gets impressive results within hours. The demo works. Everyone is excited.

Then an engineer asks a question the notebook never anticipated.

Can we trace which document this answer came from?

What happens when the vector database is down?

How do we switch from Ollama to OpenAI without rewriting the retrieval code?

Which version of the prompt produced this answer last Tuesday?

Can we prove to an auditor that the system didn't hallucinate this claim?

These are not unreasonable questions. They are exactly the questions enterprise software must answer. But they are invisible in a notebook. A notebook has one happy path. Enterprise software needs every path — especially the unhappy ones — handled deliberately.

This is the gap between an AI demo and an AI platform.

The Three Gaps

Three specific gaps separate demos from platforms.

The observability gap. A notebook produces output. A platform produces output *and* an audit trail. Which model ran? Which prompt version was used? Which documents were retrieved? How confident is the answer? Without this metadata, AI output is indistinguishable from guesswork.

The provider coupling gap. A notebook calls `openai.ChatCompletion.create()`. A platform calls `chatProvider.complete(prompt)`. The difference is indirection — and indirection is what makes provider swaps possible without rewriting business logic.

The degradation gap. A notebook crashes when an API is unreachable. A platform degrades gracefully — keyword search continues when vectors are unavailable, enrichment falls back to regex when the LLM is down, retrieval works without the knowledge graph. Every external dependency is optional by design.

Concern	Demo	Platform
Provider coupling	Hardcoded API calls	Abstract interfaces with conditional activation
Failure handling	Exceptions propagate to user	Graceful degradation with fallback implementations
Prompt management	Strings in source code	Versioned registry with audit trails
Answer quality	Subjective assessment	Automated evaluation: grounding, faithfulness, hallucination

Audit trail	None	Full explainability metadata on every inference
Extensibility	Rewrite code	Implement SPI + add configuration

This platform closes those three gaps. It is a working reference implementation — not a product, not a framework. When we set out to build it, the goal was not to create the most feature-rich AI application. It was to create a foundation where architectural concerns — modularity, degradation, explainability, auditability — are first-class, not retrofitted.

The architecture embodies a simple thesis: **AI is infrastructure, not a feature.** Treat LLMs, embedding models, vector databases, and knowledge graphs the way you treat PostgreSQL. Abstract them behind interfaces. Select them via configuration. Swap them without touching business logic.

That thesis shapes every decision that follows.

Design Rule — Business logic must never depend on a specific AI provider. Every provider integration lives behind an interface. Every interface has a fallback implementation for when the provider is unavailable. No exception.

2. The Principles

Eight principles keep hundreds of design decisions coherent. When a proposal conflicts with a principle, the proposal is rejected — or the principle is amended explicitly, with documented rationale. There is no silent drift.

AI Is Infrastructure

LLMs, embedding models, vector databases, knowledge graphs — these are infrastructure components. They belong behind interfaces. Business logic depends on abstractions. Implementations activate based on configuration.

A document enrichment service never calls a specific LLM provider. It calls an interface. Whether that interface is backed by Ollama, OpenAI, or a regex fallback is not a business concern. This pattern repeats everywhere — retrieval, embedding, reranking, evaluation.

Graceful Degradation, Not Hard Failures

In production, external services fail. We learned this the hard way on previous projects — a vector database going down at 2 AM should not take the entire search system with it. Designing for degradation from day one means every external dependency is optional.

The platform starts with only PostgreSQL. Add Qdrant: vector search activates. Add Neo4j: the graph populates. Remove any: the platform continues, reduced capabilities logged clearly, never silently.

Engineering Observation — Degradation is not a feature you add later. It is a design constraint you build against from the first line of code. Retrofitting degradation into a system that assumes all services are always available is exponentially harder than designing for it upfront.

Explainability by Default

If this platform generates an answer, it can explain itself — which intent was classified, which strategy was selected, which prompt template and version was used, which provider and model ran inference, how many chunks and graph nodes were retrieved, and what evaluation scores resulted. All generated as a byproduct of execution. No separate "explain this" pass. Zero cost in the happy path.

Design Rule — Explainability must have zero marginal cost. If explaining an answer requires additional computation, it will be disabled under load precisely when it is most needed. Build it into the execution path or don't

build it at all.

Documents Are the Knowledge Source

Users upload documents. The platform builds the knowledge. There is no manual knowledge base. The enrichment engine extracts entities, concepts, and relationships automatically during ingestion. The knowledge graph is auto-generated.

An earlier version of this project included a manual knowledge base where users could create and curate structured knowledge entries. We removed it. The realization was simple: maintaining a parallel knowledge structure alongside the document corpus created consistency problems. Which was authoritative — the uploaded document or the curated entry? What happened when the document was updated but the entry was not? The document corpus is the only source of truth. Everything else is derived.

Domain Independence

The core contains no assumptions about any industry or document type. Domain customization lives behind a pluggable interface. The included Document Intelligence application provides defaults. New domains add their own configuration — no core changes.

Testing as Architecture Documentation

Tests describe behavior, not implementation. A Staff Engineer should understand the platform by reading the test suite. Test names read like specifications.

Production Readiness

Observability, health checks, and configuration are first-class concerns designed alongside the features they observe — not afterthoughts added before release.

Modular Monolith

Module boundaries enforced at compile time. Dependency flows one way. Higher modules depend on lower modules — never the reverse.

3. Why This Platform Exists

Most AI projects begin the same way. Someone writes a script that calls an LLM API. They demo it on a laptop with one document. It works. Everyone is excited.

Then the engineering questions begin.

What happens when the vector database is down? Which prompt version generated that answer? How do we add a new provider alongside the existing one without rewriting the retrieval code? Can we trace why this answer cited that document?

These questions expose the gap between an AI demo and an AI platform. A demo needs one happy path. A platform needs every unhappy path handled deliberately.

This project closes that gap. It demonstrates how experienced engineers integrate AI into enterprise infrastructure — not as a feature bolted onto an existing system, but as infrastructure designed from first principles.

The result is a modular monolith: 9 Maven modules, a single deployable, compile-time boundaries enforcing dependency direction. It is not a product. It is not a framework. It is a foundation — code another engineering team could adopt as the starting point for their own AI application.

What's in scope: multi-provider AI orchestration, hybrid search across keyword, vector, and graph sources, automatic semantic enrichment during ingestion, GraphRAG with knowledge graph traversal, versioned prompt management, a queryable model capability database, automated evaluation of every answer, full explainability metadata, a reusable workflow engine, production metrics, and 157 automated tests.

What's deliberately absent: microservices, multi-agent systems, model fine-tuning, user management, billing. These are valid concerns that would distract from the core architecture.

Part II — Architecture

4. How the Platform Fits Together




Figure 4.1. Nine modules in three layers. Dependency flows downward. The application module depends on everything. Nothing depends on the application module.

The Modular Monolith

Three layers, nine modules, one direction of dependency.

At the bottom sits the audit module — the only leaf. Every other module depends on it for immutable audit logging with correlation IDs that thread through every request.

The middle layer holds seven platform modules. Each has one responsibility. Authentication. Document lifecycle. Hybrid search. AI orchestration. Knowledge graph persistence. Workspace management. Observability. When a module grows large, its internal package structure provides further separation without adding Maven modules. The goal is clarity, not a specific module count.

At the top, the assembly module wires everything together — REST controllers, UI templates, DTOs, and the configuration that conditionally activates provider beans based on which infrastructure is available. Future applications would create their own assembly modules on the same platform foundation.

Approach	Benefits	Costs
Modular Monolith	Compile-time boundaries; single deployable; clear dependency direction	No runtime isolation; modules share JVM
Microservices	Independent deployment; isolated failure domains; per-service scaling	Network complexity; eventual consistency; operational overhead
Single JAR (no modules)	Simplest to build	No dependency enforcement; rapid architectural decay

External Dependencies




Figure 4.2. Four external systems. PostgreSQL is required. Everything else — Qdrant, Neo4j, LLM providers — is optional.

Container Architecture — Enterprise AI Platform

C4 Level 2

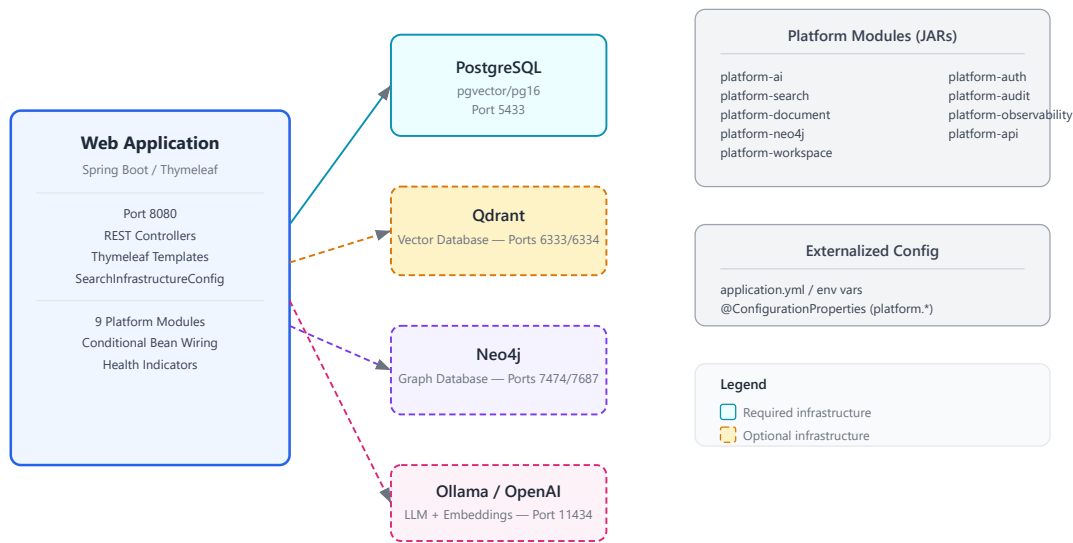


Figure 4.3. A single Spring Boot process communicates with its dependencies over standard protocols. Optional services have dashed borders — the architecture expects them to sometimes be absent.

The platform communicates with four external systems. PostgreSQL is the only hard requirement. Qdrant provides vector search. Neo4j stores the knowledge graph. Ollama or OpenAI provide LLM inference and embeddings. Each additional service adds capability but is never required.

Spring Boot 3.3 with Java 21 was chosen for ecosystem breadth and enterprise familiarity — not because it is the fastest or most lightweight framework, but because a reference implementation should prioritize clarity over novelty.

Why Java?

A reasonable question: in an era where most AI infrastructure is built in Python, why choose Java?

Ecosystem maturity. Spring Boot provides battle-tested infrastructure — security, JPA, scheduling, Actuator, property binding — that would need to be built or assembled from disparate libraries in other ecosystems. This lets the architecture focus on AI concerns rather than application infrastructure.

Compile-time safety. Nine Maven modules with directional dependencies enforced at compile time. A Python project with equivalent modularity would require runtime enforcement through import linters and convention. Java's type system and Maven's dependency resolution make module boundaries structural, not aspirational.

Concurrency model. Java 21's virtual threads enable straightforward parallel processing of document ingestion, embedding generation, and enrichment — without the complexity of reactive programming models or the overhead of platform thread pools.

Operational ecosystem. Micrometer for metrics, Prometheus for collection, Actuator for health checks — the Java observability ecosystem is mature and standardized. These are not libraries you need to evaluate and integrate. They are conventions the ecosystem has already settled.

Long-term maintainability. A reference implementation meant to be studied by other engineering teams benefits from explicit types, clear interfaces, and a language where refactoring is supported by tooling rather than test coverage alone.

None of this means Java is universally superior to Python for AI work. It means Java was the right choice *for this project* — a modular monolith prioritizing architectural clarity, long-term maintainability, and enterprise integration patterns over development speed.

5. The Document Journey

A document's journey through the platform touches text extraction, semantic enrichment, chunking, embedding generation, and three separate persistence stores — all while the user who uploaded it has moved on to other things.

The Asynchronous Choice

When we first built the ingestion pipeline, we considered processing documents synchronously during the HTTP request. The upload endpoint would extract text, chunk, embed, and index — all before returning a response. This worked on a laptop with a 2-page PDF. It failed completely with a 200-page financial report.

Document processing takes seconds to minutes. HTTP request threads should never block that long. The fix was a scheduled worker that polls for pending jobs every ten seconds. Upload creates a job and returns immediately. The worker does the heavy lifting asynchronously.

Engineering Observation — *"It works on my machine" is the most dangerous phrase in software engineering. What works for a 2-page PDF on localhost fails catastrophically for a 200-page PDF behind a load balancer. Design for the worst case, not the demo case.*



Document Ingestion Pipeline

Figure 5.1. The ingestion pipeline. Upload creates a pending job. A scheduled worker polls every 10 seconds. Extraction, enrichment, chunking, and embedding run sequentially. Results land in PostgreSQL, Qdrant, and Neo4j.

The pipeline is sequential by design, not by limitation. Parallel enrichment and embedding would improve throughput but make the code harder to trace. The reference implementation prioritizes clarity. Production deployments may add parallelism — the architecture does not prevent it.

Enrichment Before Chunking

One design choice emerged during testing: enrichment runs before chunking, not after. Our first implementation chunked first, then enriched each chunk individually. Entity extraction quality was poor — a name like "Apple" is ambiguous in a 200-word chunk but obvious in a 5,000-word financial report. Reversing the order — enrich the full document, then chunk — improved entity disambiguation significantly. This looks obvious in retrospect. It was not obvious when we wrote the first implementation.

Lesson Learned — *Context size matters for extraction quality. Chunking destroys context. Run extraction on the largest coherent text unit available, then chunk the results. This principle applies beyond document processing — it applies to any system that extracts structured information from unstructured text.*

Degradation in Practice

Each step in the pipeline degrades independently. No embedding provider configured? Chunks are stored keyword-only, and search still works. No LLM available? Regex patterns extract entities with reasonable accuracy. No Neo4j? Enrichment still runs — only graph persistence is skipped. The pipeline is designed as a chain of independent, fallible steps, not a monolithic transaction.

6. How AI Is Wired In

Answering a question with AI is not one operation. It is twelve — and each produces metadata that feeds the explainability trail.




Figure 6.1. The twelve-step inference pipeline. Every step produces metadata. Validation, grounding, and evaluation run even when earlier steps produce warnings — the platform surfaces concerns rather than hiding them.

The Full Pipeline

Intent classification. Before launching expensive retrieval, the system classifies the query. An index inspection routes to keyword-only. A substantive question enters the full RAG pipeline. This classification is deterministic — same query, same path, every time.

Strategy selection. The classified intent maps to a retrieval strategy. Keyword for exact lookups. Hybrid for complex questions. Semantic for conceptual exploration. The mapping is a simple lookup table, not a machine learning model — explainable, auditable, and sufficient for current query types.

Retrieval. Search fans out to keyword, vector, and optionally graph sources. Results are fused through weighted linear combination and reranked.

Context assembly and prompt building. Retrieved chunks and analysis objectives are assembled. The prompt registry resolves the latest template. Variables are substituted.

Provider selection. The router selects which LLM backend to use based on model prefix, capability requirements, preference, and availability.

LLM inference, validation, grounding. The selected provider generates raw answer text. Temporal consistency is checked. Claims are validated against available sources. Sources are re-attributed with confidence scores.

Evaluation and explainability. Every decision from the pipeline is captured — traceable, auditable, deterministic. The `explain()` method produces a human-readable summary.

The Provider Architecture




Figure 6.2. Five service provider interfaces decouple business logic from infrastructure. Implementations activate conditionally. Fallbacks ensure the platform works without any optional service.

Approach	Example	Consequence
Provider SPI	<code>chatProvider.complete(prompt)</code>	Swap providers without changing orchestration
Direct API calls	<code>openaiClient.chat(prompt)</code>	Every provider change requires rewriting business logic
Factory pattern	<code>ProviderFactory.get("openai")</code>	Centralized but still couples factory to provider names

The provider architecture is the Strategy pattern applied to AI infrastructure. Business logic depends on interfaces. Implementations activate based on configuration. The two never meet in source code. This is not novel — but applying it consistently to AI, where most projects hardcode provider specifics, is what makes the architecture worth studying.

Part III — Subsystems

7. Semantic Enrichment

Keyword search finds exact matches. Vector search finds semantically similar passages. Neither understands that "Acme Corporation" is an organization, or that "\$45.2 million" relates to "revenue," or that "Dr. Sarah Chen" signed a document on a specific date.

This structured understanding — who, what, when, how much — is what enables intelligent retrieval. Without it, a query like "What contracts did Acme sign in Q2?" relies on the hope that relevant documents contain those exact words near each other.

Approach	How it works	Limitation
Keyword search	Exact token matching	Misses synonyms, entities, relationships
Vector search	Semantic embedding similarity	Misses structured relationships (entities, concepts)
Semantic enrichment	Extract entities, concepts, relationships	Requires extraction accuracy; builds structured knowledge



Figure 7.1. Enrichment runs after text extraction, before chunking. Full document context improves entity disambiguation. Results flow to the knowledge graph when available.

Dual-Mode Extraction

The enrichment engine uses two modes, selected automatically: LLM-based extraction when a chat provider is available (higher accuracy, understands context), and regex pattern matching when no LLM is available (catches organizations, persons, dates, monetary amounts — enough for basic retrieval enhancement).

Mode	When used	Accuracy	Dependency
LLM-based	Chat provider available	Higher — understands context	Requires LLM
Regex	No LLM available	Baseline — patterns only	No dependency

Why Enrich During Ingestion

Enrichment is expensive — LLM calls, regex passes over full documents. Query-time latency budgets are measured in milliseconds. Pre-computing enrichment during ingestion means retrieval is fast and enrichment quality is known before any question is asked. This is an architectural trade-off: ingestion latency for query-time speed.

Provenance

Design Rationale: Every graph node carries provenance. Without it, "Acme Corporation is an ORGANIZATION" could be a hallucination. With provenance, you trace it to the exact document, chunk, model, prompt version, and timestamp. The knowledge graph is auditable — not in theory, in every node, by construction.

8. GraphRAG

Traditional RAG retrieves chunks by keyword and vector similarity. It works when relevant documents share vocabulary or embedding proximity with the query. It fails when documents are semantically related but textually distant.

Example: a contract signed by Acme Corporation and a financial report mentioning Acme's Q2 revenue share no keywords, and their embeddings may sit far apart in vector space. But they are undeniably related — through the entity "Acme Corporation."

Approach	Strength	Weakness
Keyword retrieval	Exact matches; fast	Misses synonyms, conceptual connections
Vector retrieval	Semantic similarity	Misses structured entity relationships
Graph traversal	Entity/concept connections	Requires populated knowledge graph

A knowledge graph captures these relationships explicitly. Traversing it reveals connections that keyword and vector search miss. This was the motivation for GraphRAG — not theoretical elegance, but a real retrieval blind spot.




Figure 8.1. Three retrieval sources converge in weighted fusion. Graph results boost existing candidates — they don't compete with them. When the graph is unavailable, retrieval degrades to keyword + vector.

The Central Design Choice

Graph results **boost, not replace**. A graph traversal finding that adds up to 15% score increase for candidates with related graph nodes. A document at position #15 by keyword + vector might move to #5 because it mentions the same entities as a top-ranked document. Graph context nudges rankings rather than overriding them.

A Cautionary Tale

An honest admission: GraphRAG was documented as a feature for months before it was implemented. The architecture manual described knowledge graph participation in retrieval. The code used a permanent no-op fallback. Building the adapter that bridged enrichment output to the retrieval SPI closed that gap. The lesson: document what *is* built, not what *should be* built. Documentation ahead of implementation creates confusion. Documentation trailing implementation is always accurate.

9. The Retrieval Decision Layer

Running every retriever for every query wastes resources. An index inspection does not need vector search. A conceptual question benefits from semantic similarity. A complex analytical question benefits from hybrid search with graph augmentation.

The solution: a decision layer whose sole responsibility is choosing which retrievers to use based on classified intent.




Figure 9.1. Intent drives strategy deterministically. Every routing decision is logged for auditability.

Query Type	Strategy	Why
Index inspection, corpus discovery	Keyword only	Exact match queries; vector search adds no value
Question answering, analysis	Hybrid	Complex queries benefit from keyword + vector fusion
Document research, lookup	Semantic	Conceptual queries match better on meaning than keywords

Deterministic Over ML-Based

The intent classifier uses keyword rules, not machine learning. We made this choice deliberately. Deterministic routing means the same query always takes the same path — debuggable, auditable, predictable. A trained classifier would likely be more accurate, but we valued transparency over accuracy for this component.

The consequence is worth stating plainly: misclassification propagates deterministically. A query classified as keyword-only when it needs hybrid search will produce poorer results every single time. We accepted this. If classification accuracy becomes critical, the classifier can be replaced without touching the orchestrator — the interface supports it.

Design Rule — Make components individually replaceable. The classifier can be upgraded from keyword rules to an ML model without any changes to the retrieval orchestrator, the prompt builder, or the inference pipeline. Each component owns its interface. No component owns another component's implementation.

The Fusion Algorithm

Weighted linear combination — not Reciprocal Rank Fusion. RRF requires rank positions from every retriever, which are not always available (graph search returns discovery results, not ranked lists). Weighted linear fusion works with raw scores and is simpler to explain. The algorithm is accurately named — we don't claim to implement RRF.

```
score = (keyword×0.40 + vector×0.40 + confidence×0.20) × docTypeWeight + graphBoost
```

10. Prompts as Engineering Assets

A prompt is the interface between your application and an LLM. Change the prompt, and answers change — sometimes subtly, sometimes dramatically. In most AI projects, prompts are strings embedded in source code. Changed without versioning. Untraceable after deployment. Impossible to A/B test.

This works for prototypes. It fails for platforms where the prompt *is* the product.

Engineering Observation — *Prompt versioning appears unnecessary until the first prompt regression. A small wording change degrades answer quality, nobody can trace which answers used the old version, and rollback means reverting a git commit and redeploying. After that experience, versioning becomes mandatory.*

Consider what happens without a registry. A team member changes a prompt to fix a specific issue. The change ships. A week later, someone notices that answers to a different class of questions have degraded. Which prompt version produced those degraded answers? Nobody knows. The prompts were strings in code — they changed, and the only record is a git diff buried in a commit with an unhelpful message.

A versioned registry makes this problem disappear. Every prompt has a qualified ID — `rag-answer/v1` , `entity-extraction/v1` . Every inference records which prompt version was used. When answer quality changes, you can correlate it to a prompt change immediately. You can run the same query against v1 and v2 and compare evaluation scores. Prompt changes become auditable, testable, and reversible.

Nine categories organize prompts by purpose — retrieval, summarization, extraction, classification, evaluation, reasoning, workflow, graph, search, system. Categories are not just organizational. They enable discovery:

`findByCategory(RETRIEVAL)` returns all prompts designed for retrieval tasks. A new team member can browse the registry and understand what prompts exist and how they are used.

The registry is deliberately simple. No template inheritance — prompts are independent. No dependency between prompts — changing one never affects another. No runtime compilation — substitution is simple string replacement. Complexity should be justified by need, not added preemptively.

Design Rule — *Every prompt must be versioned. Every inference must record which prompt version was used. This is not optional. It is the minimum requirement for auditing, debugging, and improving prompt quality over time.*

11. Models as Infrastructure

Different models have different capabilities. GPT-4o supports vision and structured output. nomic-embed-text only does embeddings. qwen2.5:7b is fast and local but lacks vision. Business logic must not hardcode assumptions about model capabilities.

The model capability registry describes every known model across eight capability dimensions plus quantitative metadata: context window, output tokens, latency estimates, cost estimates, temperature, use cases.

Selection Strategy	How it works	When it applies
Model prefix	<code>openai:gpt-4o</code> routes to openai provider	Explicit model choice
Capability match	Request STREAMING → find capable provider	Capability requirement
Preferred provider	Caller preference	Known provider preference
First available	First healthy provider	Fallback

The router receives all available implementations via dependency injection. Adding a provider means implementing the interface and adding a Spring component — the router discovers it without code changes. This is the Strategy pattern applied to AI infrastructure.

The architecture prevents business logic from selecting providers directly. Services request capabilities — streaming, embeddings, JSON output — and the router maps capabilities to providers. This indirection is what makes adding providers a zero-change operation for orchestration code.

12. Workflows Without Hardcoding

Document intelligence involves multi-step processes. Upload leads to extraction, enrichment, chunking, indexing, review, completion. Hardcoding these steps in a controller couples process logic to HTTP concerns and prevents reuse.

Approach	Flexibility	Complexity	When to use
Controller logic	Low — hardcoded	Low	Single fixed workflow
Workflow engine	High — configurable	Medium	Multiple workflows, reusable
BPMN engine	Very high	High	Complex branching, enterprise BPM

The workflow engine separates process definition from execution. Five methods cover the API. Two workflows come pre-registered. Steps declare handler types as extension points.

The engine deliberately avoids BPMN frameworks or state machine libraries. Five methods on an interface. Deterministic linear transitions. Not the most powerful workflow engine — the simplest one that demonstrates the concept. A production system might adopt Camunda or Temporal, and the interface was designed so that substitution would not require changing callers.

Part IV — Operations

13. The Audit Trail

If this platform generates an answer, it can explain itself. Not because someone added a feature — because the orchestration layer records every decision as it happens. Intent classification. Strategy selection. Prompt template and version. Provider and model. Result counts from every retriever. Fusion method. Reranking status. Evaluation scores.

The metadata is generated during normal execution. There is no separate "explain this" pass, no post-hoc analysis, no second LLM call. Every code path produces metadata by construction.

The design principle: explainability must have zero marginal cost. If explaining an answer requires additional computation, it will be disabled under load. By making explanation a byproduct of execution, it stays on by default, in every environment, forever.

14. Observing What Happens

A platform running AI workloads needs metrics designed for AI workloads — not just request counts and error rates, but inference latencies broken down by provider and model, embedding batch durations, retrieval timings by strategy, enrichment entity counts, evaluation score distributions.

The observability module provides Micrometer metrics for every AI operation. Timers with percentile histograms. Gauges for provider availability. Summaries for evaluation scores. All exported to Prometheus and surfaced through health checks.

The metric definitions stabilize first. Full integration follows. This ordering is deliberate: define the interface, validate with a few integrations, then roll out broadly. Wiring everything at once produces metrics nobody asked for.

15. Security

The security model is standard by design. JWT with BCrypt-12 hashing. Refresh tokens hashed before storage, rotated on use. Both stateless API authentication and session-based form login. Provider API keys injected via environment variables — never committed. Input enters through variable placeholders in vetted prompt templates. Responses are HTML-escaped.

Design Rule — *Innovation in security is a bug, not a feature. Use proven patterns. Do not invent authentication schemes. Do not roll your own cryptography. The platform's security is deliberately unremarkable because remarkable security is usually broken security.*

16. How We Test This




Figure 16.1. Five testing layers. Default build runs everything except browser tests, which use a separate profile to keep the primary build fast.

Layer	Count	Validates
Unit	44	Individual components — registries, router, evaluation
Integration	91	Subsystem interaction with real Spring context
Architecture	10	End-to-end behavior — enrichment, retrieval, evaluation
Resilience	9	Graceful degradation — fallback, unavailable services
Contract	5	SPI stability — every provider satisfies abstract contract
Browser	12	User journeys via Playwright (separate profile)

The test suite is structured as executable architecture documentation. Tests describe behavior: "routes openai:gpt-4o to openai provider" — not "testProviderRouter." A Staff Engineer should understand the platform by reading the tests.

17. Deployment

One command starts the platform. One configuration file controls it. Every custom property lives under the `platform.*` namespace. Every external service is optional — the platform starts with PostgreSQL alone.




Figure 17.1. The knowledge graph database lives in a separate Docker profile. Every additional service adds capability but is never required. A reference implementation should be trivial to run on a development machine.

Part V — Reflection

The preceding chapters described how the platform works. The remaining chapters step back and ask: what did we learn? What rules emerged? What trade-offs proved correct? What would we do differently?

These questions separate documentation from architecture. Documentation describes what was built. Architecture explains why it was built that way — and what the builder learned in the process.

18. Extending the Architecture

The architecture was designed for extension. Four patterns cover the most common needs.

Adding a provider. Implement the chat provider interface. Add component annotation with conditional property activation. Register model capabilities. The router discovers it automatically.

Adding a domain. Implement the domain configuration interface. Provide it as a Spring bean. Concepts, objectives, hierarchies, and instructions become available without core changes.

Adding a workflow. Define steps and transitions. Call the workflow engine's start method. The engine handles state.

Adding a prompt. Register a template. It is immediately versioned, categorized, and discoverable.

The goal was never to predict every extension point. It was to establish patterns — interface, conditional activation, registry, dependency injection — that make extension feel natural rather than like working around constraints.

19. Design Rules

After building this platform, certain rules emerged that now guide every change. They are not theoretical. Each one was learned the hard way — by breaking it, experiencing the consequences, and fixing the result.

Business logic never depends on providers. Every AI capability — chat completion, embedding, reranking, enrichment — lives behind an interface. Orchestration code calls the interface. Implementations activate based on configuration. This is the single most important rule. It is what makes provider swaps possible without touching business logic.

Documents are the single source of truth. Knowledge is extracted from documents during ingestion. There is no manual knowledge base. If a fact appears in the knowledge graph, it came from a document. If it did not come from a document, it is not in the graph. This rule eliminated an entire class of consistency problems.

Every prompt is versioned. Prompts change. When a prompt changes, answers change. Without versioning, you cannot compare answers before and after a prompt change. Without versioning, you cannot audit which prompt produced which answer. The prompt registry makes versioning automatic.

Every answer is explainable. The orchestration layer records every decision — intent, strategy, prompt, provider, model, results, scores — as a byproduct of execution. There is no separate explanation step. Explanation has zero marginal cost.

Every optional dependency degrades gracefully. The platform starts with PostgreSQL alone. Add Qdrant: vector search activates. Add Neo4j: the graph populates. Remove any: the platform continues. Conditional bean activation, fallback implementations, and availability checks implement this rule everywhere.

Every extension begins with an interface. If a component might need to be replaced — and in AI, everything might need to be replaced — define an interface first. Implementations come second. This is not premature abstraction. It is acknowledging that AI infrastructure evolves faster than application code.

Infrastructure is replaceable. Qdrant today might be Weaviate tomorrow. Ollama today might be vLLM tomorrow. The architecture assumes this and designs for it. Business logic never imports a specific vector database client.

Retrieval should be deterministic before it becomes intelligent. Intent classification uses keyword rules, not ML. Strategy selection is a lookup table. Fusion is weighted linear combination. Every step is deterministic, auditable, and debuggable. Intelligence can be added later — but a system that is not deterministic cannot be debugged, and a system that cannot be debugged cannot be trusted.

Wiring must be verified, not assumed. In a Spring application with conditional beans and optional dependencies, it is possible for an entire pipeline to compile, pass tests, and never execute. Wiring verification tests — tests that confirm beans are injected and called — prevent this. This rule was added after we discovered the AI inference pipeline was dead code.

These rules are not permanent. They will evolve as the platform evolves. But they represent hard-won engineering judgment. Breaking any of them should require explicit justification.

20. Performance Notes

This is a reference implementation, not a production deployment. It has not been benchmarked. What follows are architectural observations.

Chunking scales linearly with document length. Embedding generation is sequential — parallelizable via virtual threads. Weighted fusion is linear in candidate count. Graph traversal is bounded at 2 hops. LLM inference timeout is configurable.

No caching layer exists. Each retrieval re-executes. This is appropriate for demonstrating architecture. It would be the first optimization in production.

21. The Trade-offs We Accepted

These decisions reflect engineering judgment, not universal truth. Each was debated, documented, and deliberately accepted.

Decision	Chosen	Alternative	Why
Architecture	Modular monolith	Microservices	Compile-time boundaries sufficient; extraction possible
Framework	Spring Boot 3.3	Quarkus	Ecosystem breadth; enterprise familiarity
Graph database	Neo4j property graph	RDF / relational	Traversal queries; schema flexibility
Fusion algorithm	Weighted linear	Reciprocal Rank Fusion	Heterogeneous scores; simpler explanation
Prompt management	Versioned registry	Inline strings	Audit trails; regression testing

Enrichment	Dual-mode (LLM + regex)	LLM only	Works without AI infrastructure
Provider routing	Deterministic 4-tier	ML-based	Explainable; auditable
UI testing	Playwright	Selenium	Modern API; auto-waits; reliability

Each decision has a dedicated ADR in the companion volume with deeper rationale, alternatives considered, and future evolution paths.

22. What We Learned

Clean interfaces make domain-specific code easy to remove. The original codebase contained 55 German BGB paragraph references and 25+ legal keywords embedded in enrichment and analysis services. They were removed without architectural changes because the interfaces had clean separation from domain logic. The lesson generalizes: invest in interface design early. It pays off when requirements shift. This is not theoretical — it happened, and the interfaces saved months of refactoring.

Unused dependencies accumulate invisibly. Five JARs from a Spring AI integration sat on the classpath for months with zero imports anywhere in the codebase. A monthly dependency analysis run would catch this — it costs nothing and prevents deadweight from accumulating. The build audit that discovered this also found Flyway dependencies (migrations written but never enabled) and transitive JARs pulled in by unused direct dependencies.

Bean wiring must be verified, not assumed. The entire AI inference pipeline — twelve steps, a dozen collaborating services — compiled cleanly and passed every test. But the pipeline was never connected to the controller. The AI page showed retrieval results without generating LLM answers. A simple wiring verification test caught this. In Spring applications, test that your beans are actually injected and called — not just that they compile.

Documentation must trail implementation, not lead it. GraphRAG appeared in the architecture manual months before the code implemented it. The knowledge graph was populated but never queried during retrieval. A permanent no-op fallback sat where real graph traversal should have been. Building the adapter closed that gap. Now documentation reflects code, not aspiration.

Prompt versioning is table stakes for any AI platform. What started as a simple registry — store templates, substitute variables — became the foundation for deterministic behavior validation, A/B testing prompts, audit trails, and regression testing. Every inference records exactly which prompt version produced it. This was not planned from the start. It became obvious the first time a prompt change caused a behavior regression and nobody could trace which answers used the old version.

Tests as documentation works. Evolving from wiring-verification tests to behavior-driven tests made the suite readable by architects who never touched the implementation. "Routes openai:gpt-4o to openai provider" communicates more than "testProviderRouter" — not because the test implementation differs, but because the name describes what the system does rather than what class it tests.

23. The Architecture Endures

In five years, the models running this platform will be different. In three years, the vector database may have been replaced. In two years, the LLM provider you use today may no longer exist. The specific choices documented in this handbook — Ollama for local inference, Qdrant for vector storage, Neo4j for graph traversal — will look dated sooner than any of us expect.

But the principles that shaped those choices will not date.

Treating AI as infrastructure rather than a feature means the interface between business logic and AI backends remains stable even as every backend changes. Designing for graceful degradation means the platform is never more fragile than its most unreliable dependency. Making explainability a byproduct of execution means every answer is auditable forever, at zero additional cost. Extracting knowledge from documents during ingestion means enrichment quality is verified before any question is asked.

These principles are independent of any specific technology. They apply whether inference runs on Ollama, OpenAI, or a model that has not been invented yet. They apply whether vectors live in Qdrant, pgvector, or a database that launches next year. They apply whether the graph lives in Neo4j, a property graph embedded in the application, or a graph database built on an entirely different paradigm.

Good software architecture does not predict the future. It makes the present coherent and the future possible. It does not bet on which technology will win. It designs interfaces that make technology choices replaceable. It does not assume services will always be available. It designs for their absence. It does not hide decisions behind opaque implementations. It surfaces them as metadata.

The architecture documented in this handbook is a snapshot — engineering judgment applied to a specific set of technologies at a specific point in time. The implementations will be replaced. The principles — abstraction, degradation, explainability, provenance, domain independence — will outlast every one of them.

A codebase is read by machines. An architecture is read by engineers. The machines will move on to newer code. The engineers will carry the principles forward.

That is what makes this worth studying.

Appendix A — Key Decisions

Twenty Architecture Decision Records document every major engineering decision. The complete volume — with introduction, decision index, architecture timeline, and cross-references — is available as [Architecture-Decision-Records.pdf](#).

ADR	Title	Status	Category
001	Modular Monolith	Accepted	Architecture
002	Spring Boot 3.3	Accepted	Technology
003	Provider Abstraction via SPI	Accepted	AI Infrastructure
004	Provider Router	Accepted	AI Infrastructure
005	Prompt Registry	Accepted	AI Infrastructure
006	Model Capability Registry	Accepted	AI Infrastructure
007	Semantic Enrichment Engine	Accepted	Knowledge
008	GraphRAG	Accepted	Knowledge
009	Neo4j as Knowledge Graph Persistence	Accepted	Knowledge
010	Qdrant as Vector Database	Accepted	Knowledge
011	Retrieval Orchestration	Accepted	Knowledge

012	Explainability by Default	Accepted	Quality
013	Evaluation Engine	Accepted	Quality
014	Workflow Engine	Accepted	Extensibility
015	AI Observability	Accepted	Operations
016	Graceful Degradation	Accepted	Operations
017	DomainConfiguration	Accepted	Extensibility
018	Provenance-Aware Knowledge Graph	Accepted	Knowledge
019	Multi-Layer Testing Strategy	Accepted	Quality
020	Design Philosophy	Accepted	Governance

Appendix B — Glossary

Capability — A feature an AI model supports: streaming, vision, JSON mode, embeddings, tool calling.

Chat Provider — The SPI for LLM backends. Every provider implements this interface.

Chunk — A segment of document text (~1200 characters) with sentence-boundary awareness and overlap, indexed for retrieval.

Citation Coverage — Fraction of AI answer claims that cite a retrieved source.

Context Window — Maximum tokens a model can process in a single request.

Domain Configuration — Pluggable interface for domain-specific rules: concepts, objectives, hierarchies, instructions.

Embedding — A 768-dimensional vector representation of text, generated by an embedding model.

Enrichment — Automatic extraction of entities, concepts, and relationships from documents during ingestion.

Evaluation — Automated quality assessment: grounding, faithfulness, hallucination indicators, relevance.

Explainability — Metadata describing how an inference was produced, generated as a byproduct of execution.

Faithfulness — How factually accurate an answer is relative to its source documents.

Fusion — Combining retrieval results from multiple sources into a single ranked list.

GraphRAG — RAG enhanced with knowledge graph traversal. Graph results boost, not replace.

Grounding — The degree to which an answer is supported by retrieved evidence.

Hallucination — AI-generated text not supported by source documents.

Hybrid Search — Retrieval combining keyword, vector, and optionally graph search.

Inference — The process of an LLM generating text from a prompt.

Intent — The classified purpose of a query: question answering, document lookup, index inspection.

Knowledge Graph — A network of entities, concepts, and relationships extracted from documents.

Model Capability Registry — Searchable database of model features, limits, and performance estimates.

Modular Monolith — Single deployable with compile-time module boundaries.

Provenance — Metadata linking data to its origin: source document, extraction method, model, prompt version, timestamp.

Prompt Registry — Versioned store of prompt templates organized by category.

Provider — An AI backend implementing a platform SPI.

Provider Router — Component that selects which provider to use. 4-tier deterministic selection.

RAG — Retrieval-Augmented Generation. Grounding LLM responses in retrieved documents.

Registry — A queryable, versioned store of known entities (prompts, models, capabilities).

Reranking — Reordering candidates: score sort followed by LLM cross-encoder if available.

Retrieval Strategy — Which retrievers to use, selected based on classified query intent.

SPI — Service Provider Interface. A pluggable contract for extending the platform.

Structured Output — LLM responses constrained to a format like JSON.

Token — A unit of text processed by an LLM (~0.75 words in English).

Vector — A numerical representation of text meaning, 768 dimensions in this platform.

Vector Database — Database optimized for high-dimensional similarity queries.

Workflow — A configurable multi-step process with defined state transitions.

Workflow Engine — Reusable component for executing workflows.

Appendix C — References

Architecture & Engineering

- **Spring Boot 3.3 Reference** — The platform's runtime framework
- **Martin Kleppmann, *Designing Data-Intensive Applications*** — Foundational thinking on data systems
- **Martin Fowler, *Patterns of Enterprise Application Architecture*** — Patterns for modular design
- **Simon Brown, C4 Model** — Architecture visualization approach

AI & Retrieval

- **Lewis et al. (2020), *RAG for Knowledge-Intensive NLP Tasks*** — Foundational RAG paper
- **Microsoft Research (2024), *GraphRAG: From Local to Global*** — Graph-enhanced retrieval
- **Neo4j Graph Data Science** — Graph algorithms for centrality and community detection
- **Qdrant Documentation** — Vector search with payload filtering

Operations

- **Micrometer** — JVM metrics instrumentation
- **Prometheus** — Metrics collection and alerting
- **Playwright** — Cross-browser automation

Companion Volumes

- [Architecture Decision Records](#) — 20 ADRs
- [Developer Guide](#) — Build, run, extend, test, contribute